

535514: Reinforcement Learning

Lecture 10 — Model-Free Prediction

Ping-Chun Hsieh

March 24, 2024

3 Major Approaches for Model-Free Prediction

1. Monte Carlo (MC)

2. Temporal Difference: TD(0) and n -step TD

3. TD(λ) and GAE

References:

Richard Sutton and Andrew Barto, Reinforcement Learning: An Introduction, 2019

Singh and Sutton, "Reinforcement Learning with Replacing Eligibility Traces," ML 1996

Schulman et al., High-Dimensional Continuous Control Using Generalized Advantage Estimation, ICLR 2016

Monte-Carlo for Policy Evaluation

- ▶ **Recall:** Monte-Carlo policy gradient
 - ▶ Use sample return G_t in the estimate of policy gradient

$$\nabla_{\theta} V^{\pi_{\theta}}(\mu) \approx \sum_{t=0}^{\infty} \gamma^t G_t \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

- ▶ **Question:** Can we use the same idea for policy evaluation (i.e., finding $V^{\pi}(s)$ or $Q^{\pi}(s, a)$)?

Monte-Carlo (MC) Method for Policy Evaluation

To find the value function V^π under a fixed policy π :

Monte-Carlo Policy Evaluation

- ▶ For **episodic** environments ➡ sample a set of trajectories $\{\tau^{(i)}\}_{i=1}^K$ and calculate average returns $\frac{1}{K} \sum_{i=1}^K G(\tau^{(i)}) \approx V^\pi(\mu)$
- ▶ For **continuing** environments ➡ sample a set of trajectories (but with proper **truncation**) and calculate average returns as an estimate of $V^\pi(\mu)$

Is MC Policy Evaluation Useful in Practice?

Yes! MC serves as a pseudo-oracle for true $V^\pi(s)$ or $Q^\pi(s, a)$

Example: Finding the “true value functions” in the TD3 paper






Figure 1. Measuring overestimation bias in the value estimates of DDPG and our proposed method, Clipped Double Q-learning (CDQ), on MuJoCo environments over 1 million time steps.



- If the policy is *deterministic*, how many trajectories do we need?
- What if the policy is *stochastic*?


Two Variants of MC Policy Evaluation: First-Visit and Every-Visit

- ▶ A visit to s : an occurrence of a state s in an episode
- ▶ **First-visit MC**: Estimate the value of a state as the average of the returns that have followed the first visit to the state in an episode
- ▶ **Every-visit MC**: Estimate the value of a state as the average of the returns that have followed all visits to the state



Example: First visit to ★? How many visits to ★?

Example: 2-State MRP



@Start state: reward = 1

@Terminal state: reward = 0

- ▶ Consider a sample trajectory: $S \rightarrow S \rightarrow S \rightarrow S \rightarrow T$
- ▶ **Question:** First-visit MC estimate of $V(S) = ?$ 4
- ▶ **Question:** Every-visit MC estimate of $V(S) = ?$
 $(4 + 3 + 2 + 1)/4 = 2.5$
- ▶ **Question:** Which estimate is better?

First-Visit MC Policy Evaluation (Formally)

Initialize $N(s) = 0$, $G(s) = 0 \ \forall s \in S$

Loop

- Sample episode $i = s_{i,1}, a_{i,1}, r_{i,1}, s_{i,2}, a_{i,2}, r_{i,2}, \dots, s_{i,T_i}$
- Define $G_{i,t} = r_{i,t} + \gamma r_{i,t+1} + \gamma^2 r_{i,t+2} + \dots + \gamma^{T_i-t} r_{i,T_i}$ as return from time step t onwards in i th episode
- For each state s visited in episode i
 - For **first** time t that state s is visited in episode i
 - Increment counter of total first visits: $N(s) = N(s) + 1$
 - Increment total return $G(s) = G(s) + G_{i,t}$
 - Update estimate $V^\pi(s) = G(s)/N(s)$

Every-Visit MC Policy Evaluation (Formally)

Initialize $N(s) = 0$, $G(s) = 0 \ \forall s \in S$

Loop

- Sample episode $i = s_{i,1}, a_{i,1}, r_{i,1}, s_{i,2}, a_{i,2}, r_{i,2}, \dots, s_{i,T_i}$
- Define $G_{i,t} = r_{i,t} + \gamma r_{i,t+1} + \gamma^2 r_{i,t+2} + \dots + \gamma^{T_i-t} r_{i,T_i}$ as return from time step t onwards in i th episode
- For each state s visited in episode i
 - For **every** time t that state s is visited in episode i
 - Increment counter of total first visits: $N(s) = N(s) + 1$
 - Increment total return $G(s) = G(s) + G_{i,t}$
 - Update estimate $V^\pi(s) = G(s)/N(s)$

An Incremental Expression of Sample Mean

- ▶ Let $z_1, z_2, z_3 \dots$ be a sequence of real numbers
- ▶ Sample mean of z_1, \dots, z_n is denoted by \bar{z}_n

$$\begin{aligned}\bar{z}_n &:= \frac{1}{n} \sum_{k=1}^n z_k = \frac{1}{n} \left(z_n + \sum_{k=1}^{n-1} z_k \right) \\ &= \frac{1}{n} \left(z_n + (n-1) \bar{z}_{n-1} \right) \\ &= \frac{1}{n} \left(z_n + (n-1) \bar{z}_{n-1} + \bar{z}_{n-1} - \bar{z}_{n-1} \right) \\ &= \bar{z}_{n-1} + \frac{1}{n} \left(z_n - \bar{z}_{n-1} \right)\end{aligned}$$

Incremental Monte-Carlo Updates

(Alternative expression of every-visit MC)

- Update $V^\pi(s)$ **incrementally** after each episode

$s_0, a_0, r_1, \dots, s_{T-1}, a_{T-1}, r_T$

For each state s_t with sample return G_t

$$\begin{aligned} N(s_t) &\leftarrow N(s_t) + 1 \\ V(s_t) &\leftarrow V(s_t) + \frac{1}{N(s_t)}(G_t - V(s_t)) \end{aligned}$$


- In **non-stationary** environments, we may instead track the exponential moving average (i.e. forget old episodes) by

$$V(s_t) \leftarrow V(s_t) + \alpha(G_t - V(s_t))$$

Comparison of First-Visit and Every-Visit MC

1. First-visit MC provides an unbiased estimate
2. Every-visit MC provides a biased but consistent estimate

Why is First-Visit MC Unbiased?



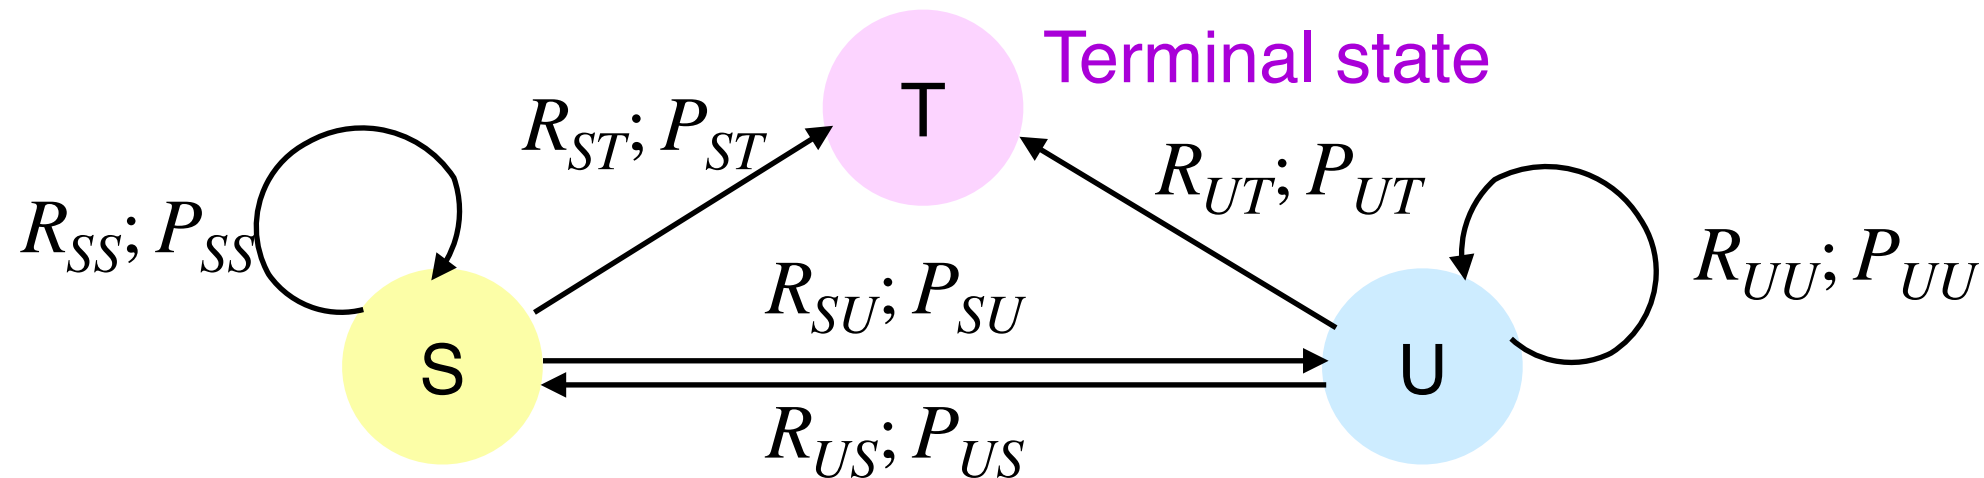
(For simplicity, suppose we use 1 trajectory τ for first-visit MC)

- In trajectory τ , suppose the first visit to state s occurs at time t
- Sample return $G_t(\tau) = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} \dots$
- Construct a first-visit MC estimate of $V^\pi(s)$ by $G_t(\tau)$


-
- **Question:** Do we have $E[G_t(\tau) | s_t = s; \pi] = V^\pi(s)$?
 - **Question:** Does this hold if we use multiple trajectories for first-visit MC?

How to Analyze Every-Visit MC? A Reduction Trick

- **Example:** Estimating $V(S)$ of a 3-state MRP (assume $\gamma = 1$)



- **Idea:** Reduce MRP to an **equivalent 2-state MRP** (in what sense?)



P_S = prob. of visiting S again before reaching T

P_T = prob. of visiting T before visiting S again

R_S = expected reward of $S \rightsquigarrow S$ transition

R_T = expected reward of $S \rightsquigarrow T$ transition

$P_S =$

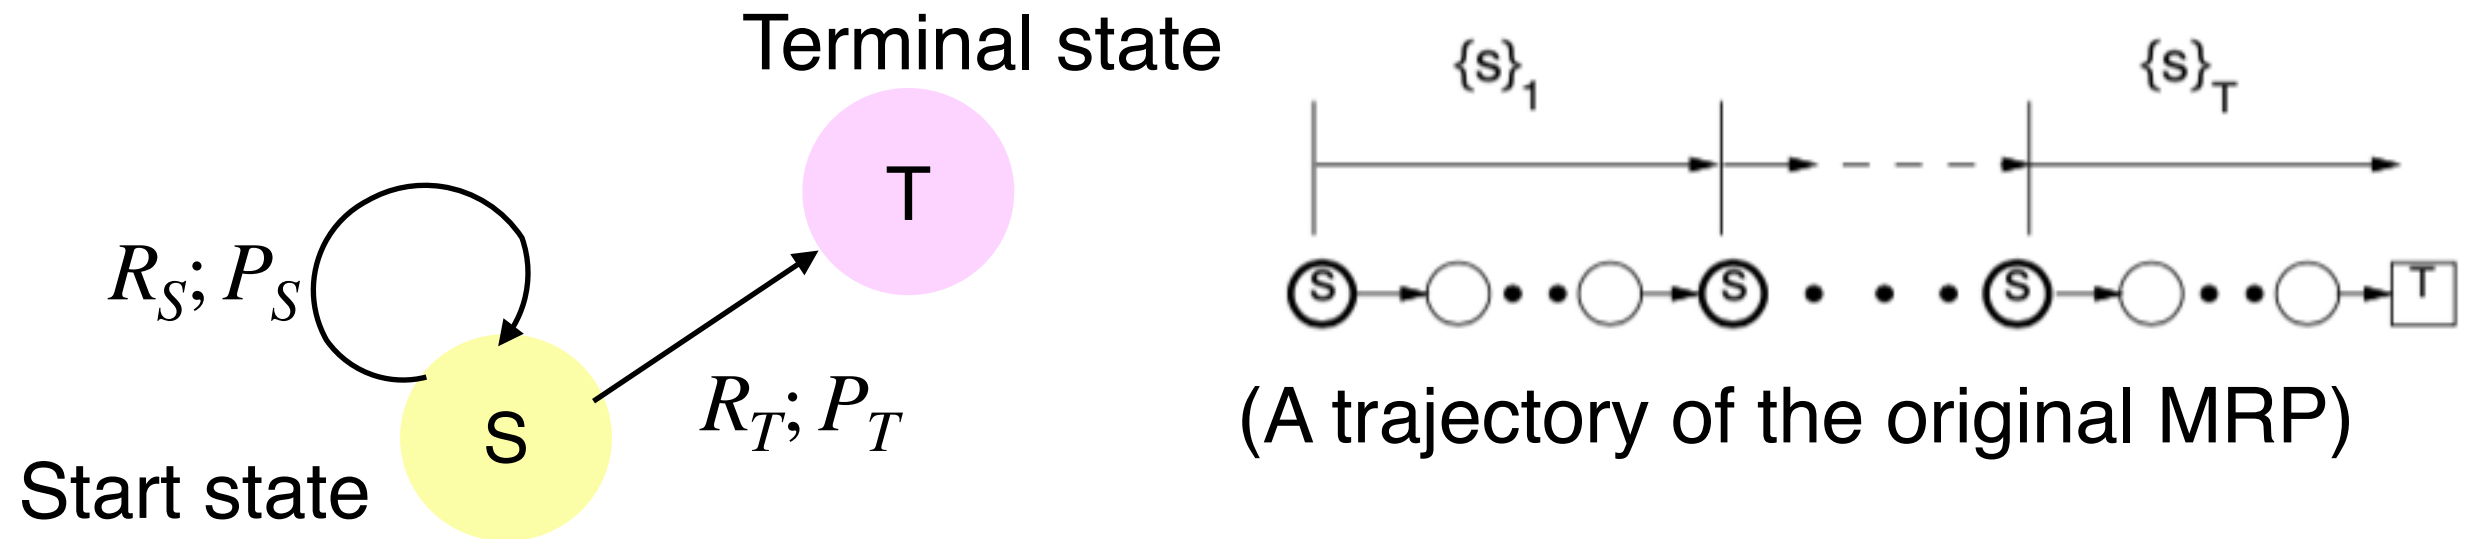
$P_T =$

$R_S =$

$R_T =$

Why is Every-Visit MC is *Biased*?

- ▶ **Next**: Estimate $V(S)$ of **any** MRP (assume $\gamma = 1$)
- ▶ **Idea**: Reduce the MRP to an equivalent 2-state MRP



Fact: True value function: $V(S) = \frac{P_S}{1 - P_S} R_S + R_T = \frac{P_S}{P_T} R_S + R_T$

- ▶ **Question**: Expected every-visit MC estimate over 1 trajectory = ?

$$\sum_{k=0}^{\infty} P_T P_S^k \left(\frac{R_S + 2R_S + \dots + kR_S + (k+1)R_T}{k+1} \right) = \frac{P_S}{2P_T} R_S + R_T$$

Why is Every-Visit MC is *Biased*? (Cont.)

- ▶ We use the same notations as in the previous page

- ▶ **Every-Visit MC is Biased but Consistent in the Limit:**

The expected every-visit MC estimate after n episodes is

$$\frac{nP_S}{(n+1)P_T}R_S + R_T$$

Thus, every-visit MC estimate is biased and the amount of bias is

$$\frac{P_S}{(n+1)P_T}R_S$$

Moreover, every-visit MC is consistent in the limit $n \rightarrow \infty$

Any Issue With Monte-Carlo Policy Evaluation?

1. MC is applicable mainly to **episodic** problems

- ▶ For continuing problems, truncation of trajectories is needed (but may incur some bias)

2. MC can only learn from complete sequences

3. MC generally has high variance

- ▶ requires a lot of samples for convergence
- ▶ might be impractical in the low-data regime

Temporal Difference (TD)

Motivating Example: A Commuter's Daily Life

- ▶ A commuter travels from NYCU back to Taichung after work

| State | Old Estimate of Time-to-Go | Sample Elapsed Time for “Today” | Predicted Total Time as of Now |
|--------------------------------|----------------------------|---------------------------------|--------------------------------|
| Leaving office | 65 | 0 | |
| Taking the bus | 55 | 8 | |
| Reaching THSR Hsinchu Station | 40 | 33 | |
| Reaching THSR Taichung Station | 10 | 64 | |
| Arriving home | 0 | 76 | |

- ▶ What technique are we using here? *Bootstrapping!*

What is TD? Comparison: TD(0) vs MC

- **Goal:** Evaluate V^π under a fixed policy π

1. Incremental every-visit Monte-Carlo: use **sample return**

$$V_{k+1}(s_t) \leftarrow V_k(s_t) + \alpha(G_t - V_k(s_t))$$

2. Temporal difference algorithm TD(0): use **estimated return**

$$V_{k+1}(s_t) \leftarrow V_k(s_t) + \alpha(r_{t+1} + \gamma V_k(s_{t+1}) - V_k(s_t))$$

- $r_{t+1} + \gamma V_k(s_{t+1})$ is the estimated return (called **TD target**)
- $r_{t+1} + \gamma V_k(s_{t+1}) - V_k(s_t) =: \delta_t$ is called the **TD error**

MC Error and TD Error

- **Fact:** MC error can be written as a sum of TD errors

$$\begin{aligned} \underbrace{G_t - V_k(s_t)}_{\text{MC error}} &= (r_{t+1} + \gamma G_{t+1}) - V_k(s_t) + \gamma V_k(s_{t+1}) - \gamma V_k(s_{t+1}) \\ &= \delta_t + \gamma (G_{t+1} - V_k(s_{t+1})) \\ &= \delta_t + \gamma \delta_{t+1} + \dots + \gamma^{T-t-1} \delta_{T-1} + \gamma^{T-t} \underbrace{(G_T - V_k(s_T))}_{=0-0} \\ &= \sum_{k=t}^{T-1} \gamma^{k-t} \delta_k \end{aligned}$$

Terminal state
↑

- **Question:** Why is the above observation intuitively useful?

Features of TD


1. TD is **model-free** (why?)

- ▶ TD learns directly from episodes **without** estimating MDP transition probabilities or reward function

2. TD learns from **incomplete** episodes by bootstrapping


- ▶ TD updates a guess towards a guess
- ▶ **Question**: Why is this a good feature (compared to MC)?

Visualization: TD(0) Backup




$$V_{k+1}(s_t) \leftarrow V_k(s_t) + \alpha \left(\textcolor{blue}{r}_{t+1} + \gamma V_k(s_{t+1}) - V_k(s_t) \right)$$

Visualization: Monte-Carlo Backup



$$V_{k+1}(s_t) \leftarrow V_k(s_t) + \alpha(\textcolor{blue}{G}_t - V_k(s_t))$$



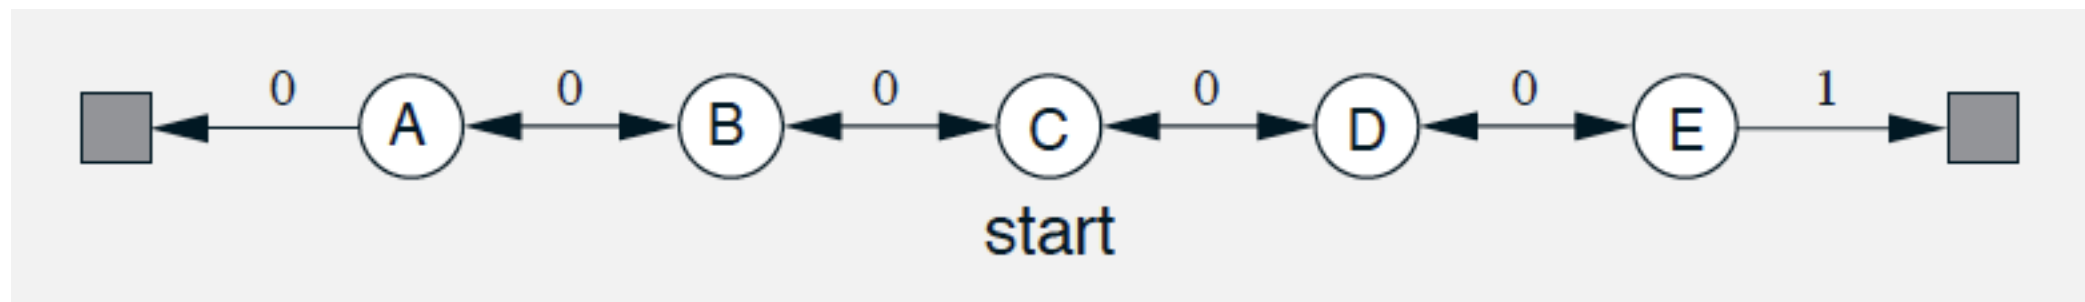
Visualization: IPE Operator Backup



$$V_{k+1}(s_t) \leftarrow \mathbb{E}_{P_\pi} [r_{t+1} + \gamma V_k(s_{t+1})]$$

Efficiency: TD(0) vs MC

- **Example:** Random walk MRP with 5 states




Extension of TD(0): n -Step TD and TD(λ)

Use n -Step Return For Prediction?

- **Recall:** update rule of TD(0)

$$V_{k+1}(s_t) \leftarrow V_k(s_t) + \alpha(r_{t+1} + \gamma V_k(s_{t+1}) - V_k(s_t))$$

- **Question:** Can we consider n steps into the future?



n -Step Bootstrapping For Prediction (Formally)

- ▶ Define the n -step estimated return

$$G_t^{(n)} = r_{t+1} + \gamma r_{t+2} + \cdots + \gamma^{n-1} r_{t+n} + \gamma^n V(s_{t+n})$$


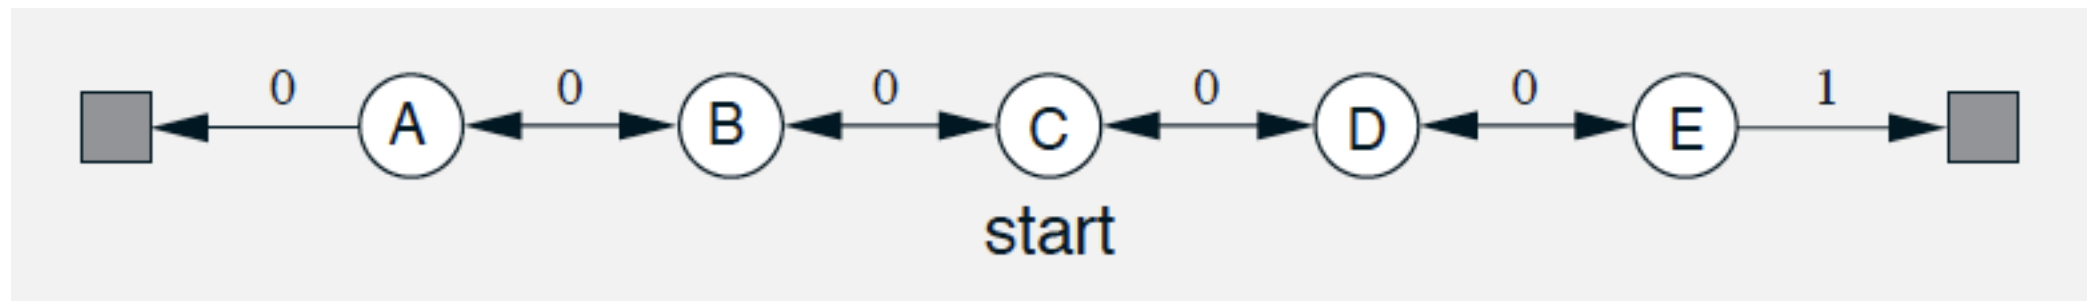
- ▶ n -step TD for policy evaluation

$$V(s_t) \leftarrow V(s_t) + \alpha \left(G_t^{(n)} - V(s_t) \right)$$


- ▶ Special case:
 - ▶ $n = 1$: standard TD(0)
 - ▶ $n = \infty$: MC

Which Value of n is Better?

- **Example:** Random walk MRP with 19 states



Combine n -Step Returns Over Different n ?



► Example:

$$G_t^{(1)} = r_{t+1} + \gamma V(s_{t+1})$$

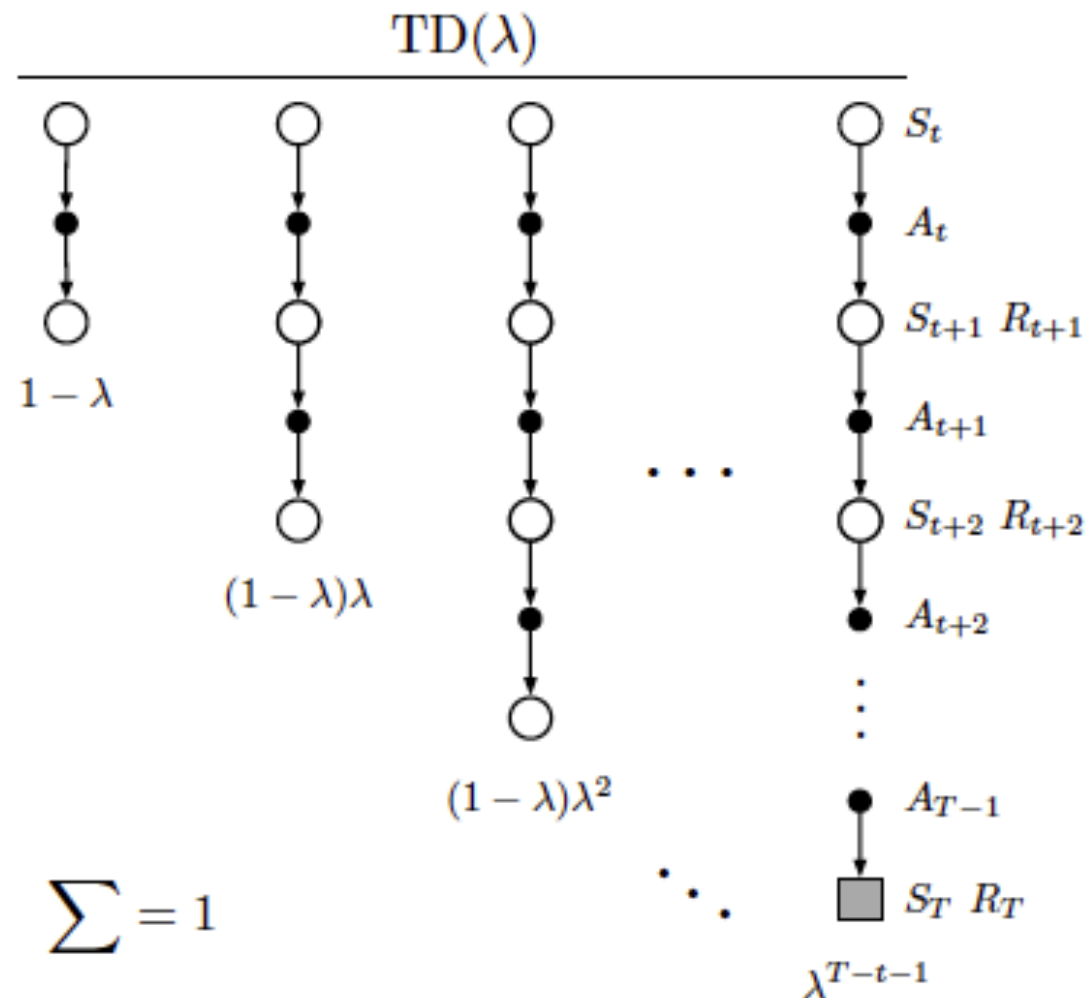
$$G_t^{(2)} = r_{t+1} + \gamma r_{t+2} + \gamma^2 V(s_{t+2})$$

$$\hat{G}_t = \frac{1}{2} \left(G_t^{(1)} + G_t^{(2)} \right)$$

$$V_{k+1}(s_t) \leftarrow V_k(s_t) + \alpha \left(\hat{G}_t - V_k(s_t) \right)$$

► Question: Any systematic way to combine n -Step Returns?

λ -Return and TD(λ)



- Define λ -return as

$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$$


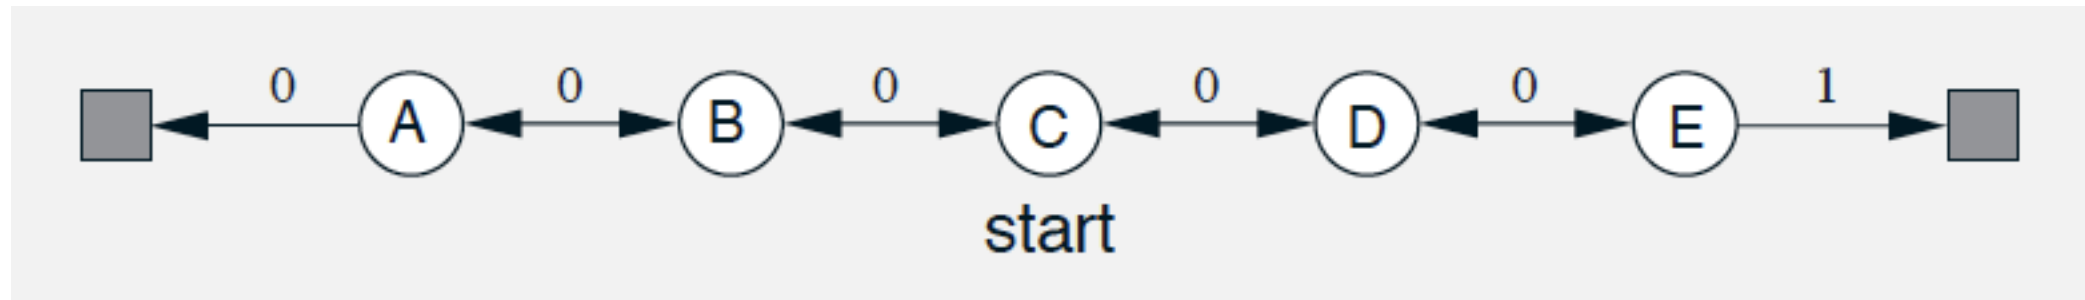
- G_t^λ combines all n -step returns using weights $(1 - \lambda)\lambda^{n-1}$

- The update rule of TD(λ) algorithm:

$$V_{k+1}(s_t) \leftarrow V_k(s_t) + \alpha(G_t^\lambda - V_k(s_t))$$

Which Value of λ is Better?

- **Example:** Random walk MRP with 19 states



Next Question: How to Estimate $A^\pi(s, a)$?

Generalized Advantage Estimator (GAE): Using TD(λ) to Estimate $A^\pi(s, a)$

Let $V(s)$ be the current estimate of true value $V^\pi(s)$

$$\hat{A}_t^{(1)} := r_{t+1} + \gamma V(s_{t+1}) - V(s_t) \quad (= \delta_t)$$

$$\hat{A}_t^{(2)} := r_{t+1} + \gamma r_{t+2} + \gamma^2 V(s_{t+2}) - V(s_t) \quad (= \delta_t + \gamma \delta_{t+1})$$

$$\hat{A}_t^{(3)} := r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \gamma^3 V(s_{t+3}) - V(s_t) \quad (= \delta_t + \gamma \delta_{t+1} + \gamma^2 \delta_{t+2})$$

$$\hat{A}_t^{(k)} := r_{t+1} + \gamma r_{t+2} + \cdots + \gamma^k V(s_{t+k}) - V(s_t) \quad (= \sum_{\ell=0}^{k-1} \gamma^\ell \delta_{t+\ell})$$

► **Fact:** $\hat{A}_t^{(\infty)} = \sum_{\ell=0}^{\infty} \gamma^\ell \delta_{t+\ell} = G_t - V(s_t)$

► **GAE Estimator:**

$$\hat{A}_t^{GAE(\gamma, \lambda)} = (1 - \gamma) (\hat{A}_t^{(1)} + \gamma \hat{A}_t^{(2)} + \gamma^2 \hat{A}_t^{(3)} + \cdots) = \sum_{\ell=0}^{\infty} (\gamma \lambda)^\ell \delta_{t+\ell}$$

Algorithm: REINFORCE With GAE

Recall: (P5) REINFORCE with advantage

$$\nabla_{\theta} V^{\pi_{\theta}}(\mu) = \mathbb{E}_{\tau \sim P_{\mu}^{\pi_{\theta}}} \left[\sum_{t=0}^{\infty} \gamma^t A^{\pi_{\theta}}(s_t, a_t) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right]$$

► REINFORCE with GAE

Step 1: Initialize θ_0 and step size η

Step 2: Sample a trajectory $\tau \sim P_{\mu}^{\pi_{\theta}}$ and make the update as

$$\theta_{k+1} = \theta_k + \eta \left(\sum_{t=0}^{\infty} \gamma^t \hat{A}_t^{GAE(\gamma, \lambda)} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right)$$

where $\hat{A}_t^{GAE(\gamma, \lambda)}$ is constructed from $V(s)$ learned by TD

(Repeat Step 2 until termination)

Some Discussions on GAE

1. Do we need to wait until the end of a trajectory to construct GAE?
2. How to efficiently calculate GAE for different t of the same trajectory?

```
def calculate_advantages(rewards, values, discount_factor, trace_decay, normalize = True):  
  
    advantages = []  
    advantage = 0  
    next_value = 0  
  
    for r, v in zip(reversed(rewards), reversed(values)):  
        td_error = r + next_value * discount_factor - v  
        advantage = td_error + advantage * discount_factor * trace_decay  
        next_value = v  
        advantages.insert(0, advantage)  
  
    advantages = torch.tensor(advantages)
```

3. Where does $V(s)$ in GAE come from?