

535514: Reinforcement Learning

Lecture 16 — TRPO and NPG

Ping-Chun Hsieh

April 18, 2024

On-Policy vs Off-Policy Methods

	Policy Optimization	Value-Based	Model-Based	Imitation-Based
On-Policy	<p>Exact PG</p> <p>REINFORCE (w/i baseline)</p> <p>A2C</p> <p>On-policy DAC</p> <p>TRPO</p> <p>Natural PG (NPG)</p> <p>PPO-KL & PPO-Clip</p>	<p>Epsilon-Greedy MC</p> <p>Sarsa</p> <p>Expected Sarsa</p>	<p>Model-Predictive Control (MPC)</p> <p>PETS</p>	<p>IRL</p> <p>GAIL</p> <p>IQ-Learn</p> <p>RLHF</p>
Off-Policy	<p>Off-policy DPG & DDPG</p> <p>Twin Delayed DDPG (TD3)</p>	<p>Q-learning</p> <p>Double Q-learning</p> <p>DQN & DDQN</p> <p>C51 / QR-DQN / IQN</p> <p>Soft Actor-Critic (SAC)</p>		


Recall: Trust Region Methods

- ▶ 3 major steps of a trust region method:

(A1) Choose a surrogate function $m_k(x)$

(A2) Specify a trust region Δ

(A3) Find an approximate optimizer in Δ



- ▶ **Question:** How to choose the surrogate function?
- ▶ **Question:** How to specify a good trust region?

Let's apply TR method and build TRPO step by step!

Goal: The Ultimate TRPO Algorithm

► Trust-Region Policy Optimization (TRPO) Algorithm:

Step 1: Initialize θ_0

Step 2: For iteration $k = 0, 1, 2, \dots$

Step 2-1: Collect trajectories by running the current policy π_{θ_k}

Step 2-2: Obtain advantage $A^{\theta_k}(s, a)$ for the current policy π_{θ_k}

Step 2-3: Update the policy by solving

$$\theta_{k+1} = \arg \max_{\theta} \mathbb{E}_{s \sim d_{\mu}^{\pi_{\theta_k}}, a \sim \pi_{\theta_k}(\cdot | s)} \left[\frac{\pi_{\theta}(a | s)}{\pi_{\theta_k}(a | s)} A^{\theta_k}(s, a) \right]$$

$$\text{subject to } \bar{D}_{KL}(\pi_{\theta_k} \| \pi_{\theta}) \leq \delta$$

Trust Region Methods for RL

- **Recall:** 3 major steps of a trust region method:

(A1) Choose a surrogate function $m_k(x)$


Approximate $d_{\mu}^{\pi_{new}}(s)$ by $d_{\mu}^{\pi_{old}}(s)$

(A2) Specify a trust region Δ

KL divergence between π_{old}, π_{new}

(A3) Find an approximate optimizer in Δ

Convexify the problem by 1st-order and 2nd-order approximation



Average Performance Difference Lemma

- ▶ **Performance Difference Lemma:**

$$V^{\pi_{new}}(\mu) - V^{\pi_{old}}(\mu) = \frac{1}{1 - \gamma} \mathbb{E}_{s' \sim d_{\mu}^{\pi_{new}}} \mathbb{E}_{a' \sim \pi_{new}(\cdot | s')} \left[A^{\pi_{old}}(s', a') \right]$$

- ▶ **Question 1:** How to interpret this result?
- ▶ **Question 2:** Under **tabular policies**, what will we have if π_{old} is obtained from π_{new} by “*one-step policy improvement*”?

To simplify notations, let's use

$$\eta(\pi_{new}) \equiv (1 - \gamma)V^{\pi_{new}}(\mu)$$
$$\eta(\pi_{old}) \equiv (1 - \gamma)V^{\pi_{old}}(\mu)$$

(A1) Surrogate Function in TRPO

Question: How about directly optimizing

$$\sum_s d_{\mu}^{\pi_{new}}(s) \sum_a \pi_{new}(a | s) A^{\pi_{old}}(s, a) = (1 - \gamma) (V_{\mu}(\pi_{new}) - V_{\mu}(\pi_{old}))$$

usually difficult to get (why?)

Approximate $d_{\mu}^{\pi_{new}}(s)$ by $d_{\mu}^{\pi_{old}}(s)$:

$$\eta(\pi_{new}) - \eta(\pi_{old}) \approx \sum_s d_{\mu}^{\pi_{old}}(s) \sum_a \pi_{new}(a | s) A^{\pi_{old}}(s, a)$$

► **Define:** Surrogate function $L_{\pi_{old}}(\pi_{new})$ in TRPO

$$L_{\pi_{old}}(\pi_{new}) = \eta(\pi_{old}) + \sum_s d_{\mu}^{\pi_{old}}(s) \sum_a \pi_{new}(a | s) A^{\pi_{old}}(s, a)$$

(A1) Why is $L_{\pi_{old}}(\pi_{new})$ a Good Surrogate Function?

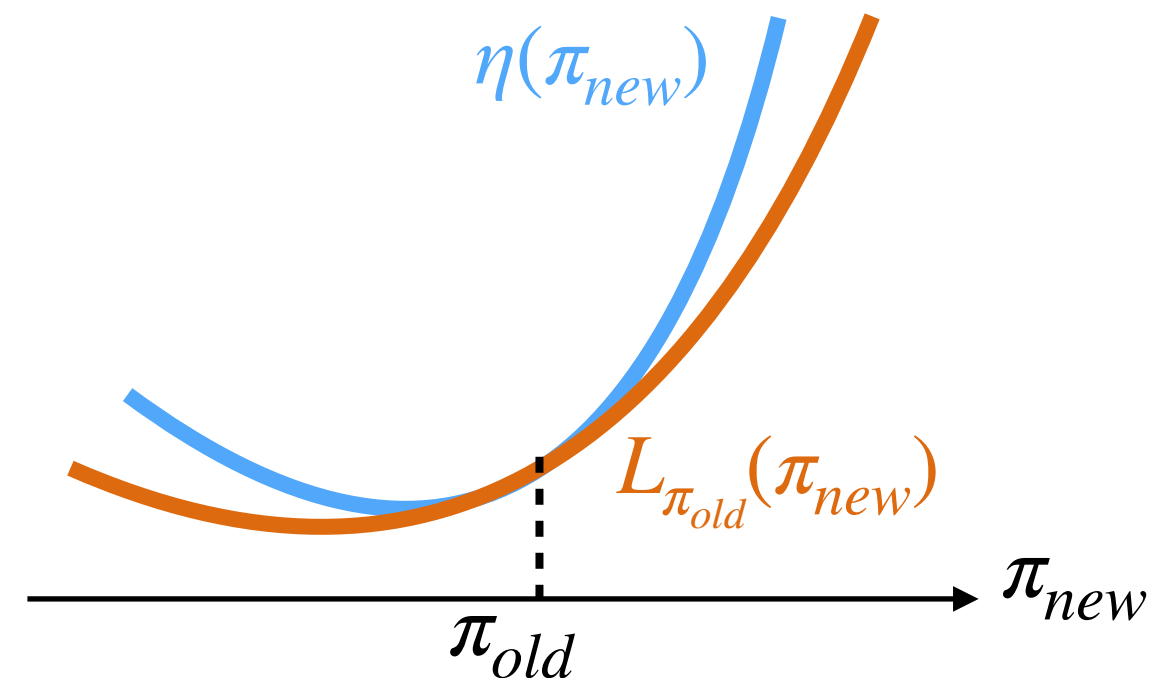
$$L_{\pi_{old}}(\pi_{new}) = \eta(\pi_{old}) + \sum_s d_{\mu}^{\pi_{old}}(s) \sum_a \pi_{new}(a | s) A^{\pi_{old}}(s, a)$$

► $L_{\pi_{old}}(\pi_{new})$ satisfy two properties: $\pi_{old} \equiv \pi_{\theta_1}$, $\pi_{new} \equiv \pi_{\theta}$

1. $L_{\pi_{\theta_1}}(\pi_{\theta_1}) = \eta(\pi_{\theta_1})$

2. $\nabla_{\theta} L_{\pi_{\theta_1}}(\pi_{\theta})|_{\theta=\theta_1} = \nabla_{\theta} \eta(\pi_{\theta})|_{\theta=\theta_1}$

(HW2 problem)



► **Intuition:** If π_{old} , π_{new} are close, then improvement in $L_{\pi_{old}}(\pi_{new})$ implies improvement in $\eta(\pi_{new})$

Kullback-Leibler divergence Between Policies

- **Notation:** Kullback-Leibler (KL) divergence

$$D_{KL}(\pi(\cdot | s) \| \tilde{\pi}(\cdot | s)) := \sum_a \pi(a | s) \log\left(\frac{\pi(a | s)}{\tilde{\pi}(a | s)}\right)$$

$$D_{KL}^{\max}(\pi \| \tilde{\pi}) := \max_s D_{KL}(\pi(\cdot | s) \| \tilde{\pi}(\cdot | s))$$


(A2) How to Specify a Trust Region Δ ?

► **Recall:** $L_{\pi_{old}}(\pi_{new})$ and $\eta(\pi_{new})$ are close if π_{old}, π_{new} are close

► **Policy Improvement Bound (PIB):** Let $\varepsilon := \max_{s,a} |A^{\pi_{old}}(s, a)|$

$$\eta(\pi_{new}) \geq \underbrace{L_{\pi_{old}}(\pi_{new}) - \frac{4\varepsilon\gamma}{(1-\gamma)^2} D_{KL}^{\max}(\pi_{old} \parallel \pi_{new})}_{=: Z_{\pi_{old}}(\pi_{new})}$$

► **Question:** How to specify a trust region?



1st attempt: $D_{KL}^{\max}(\pi_{old} \parallel \pi_{new}) \leq \delta$

Is $D_{KL}^{\max}(\pi_{old} \parallel \pi_{new})$ easy to evaluate?

(A2) How to Specify a Trust Region Δ ?

► **Policy Improvement Bound (PIB):** Let $\varepsilon := \max_{s,a} |A^{\pi_{old}}(s, a)|$

$$\eta(\pi_{new}) \geq \underbrace{L_{\pi_{old}}(\pi_{new}) - \frac{4\varepsilon\gamma}{(1-\gamma)^2} D_{KL}^{\max}(\pi_{old}||\pi_{new})}_{=: Z_{\pi_{old}}(\pi_{new})}$$

2nd attempt: Use $\bar{D}_{KL}(\pi_{old}||\pi_{new})$ instead of $D_{KL}^{\max}(\pi_{old}||\pi_{new})$

$$\bar{D}_{KL}(\pi_{old}||\pi_{new}) := \mathbb{E}_{s \sim \pi_{old}} [D_{KL}(\pi_{old}(\cdot | s) || \pi_{new}(\cdot | s))]$$

Trust region in TRPO: $\bar{D}_{KL}(\pi_{old}||\pi_{new}) \leq \delta$

Put Everything Together

► Trust-Region Policy Optimization (TRPO) Algorithm:

Step 1: Initialize θ_0

Step 2: For iteration $k = 0, 1, 2, \dots$

Step 2-1: Collect trajectories by running the current policy π_{θ_k}

Step 2-2: Obtain advantage $A^{\theta_k}(s, a)$ for the current policy π_{θ_k}

Step 2-3: Update the policy by solving

$$\theta_{k+1} = \underset{\theta}{\text{arg max}} L_{\pi_{\theta_k}}(\pi_{\theta}) \quad \left(\equiv \arg \max_{\theta} \mathbb{E}_{s \sim d_{\mu}^{\pi_{\theta_k}}, a \sim \pi_{\theta_k}(\cdot | s)} \left[\frac{\pi_{\theta}(a | s)}{\pi_{\theta_k}(a | s)} A^{\theta_k}(s, a) \right] \right)$$

subject to $\bar{D}_{KL}(\pi_{\theta_k} \| \pi_{\theta}) \leq \delta$

One remaining practical issue with TRPO...

$$\theta_{k+1} = \arg \max_{\theta} L_{\pi_{\theta_k}}(\pi_{\theta}) \quad \left(\equiv \arg \max_{\theta} \mathbb{E}_{s \sim d_{\mu}^{\pi_{\theta_k}}, a \sim \pi_{\theta_k}(\cdot | s)} \left[\frac{\pi_{\theta}(a | s)}{\pi_{\theta_k}(a | s)} A^{\theta_k}(s, a) \right] \right)$$

subject to $\bar{D}_{KL}(\pi_{\theta_k} \| \pi_{\theta}) \leq \delta$

How to efficiently solve this constrained problem?

(A3) Find an Approximate Optimizer in Δ

$$\theta_{k+1} = \arg \max_{\theta} L_{\pi_{\theta_k}}(\pi_{\theta}) \quad \left(\equiv \arg \max_{\theta} \mathbb{E}_{s \sim d_{\mu}^{\pi_{\theta_k}}, a \sim \pi_{\theta_k}(\cdot | s)} \left[\frac{\pi_{\theta}(a | s)}{\pi_{\theta_k}(a | s)} A^{\theta_k}(s, a) \right] \right)$$

subject to $\bar{D}_{KL}(\pi_{\theta_k} \| \pi_{\theta}) \leq \delta$

► Idea: Approximation

1. Linear approximation to the objective $L_{\theta_k}(\theta)$
2. Quadratic approximation to the KL constraint

► The problem under approximation:

$$\begin{aligned} &\text{Maximize} && (\theta - \theta_k)^{\top} \nabla_{\theta} L_{\theta_k}(\theta) |_{\theta=\theta_k} \\ &\text{subject to} && \frac{1}{2}(\theta - \theta_k)^{\top} H(\theta - \theta_k) \leq \delta \end{aligned}$$

Hessian of $\bar{D}_{KL}(\pi_{\theta_k} \| \pi_{\theta})$

► Question: Why is this approximation helpful?

- ▶ The problem under approximation:

$$\begin{array}{ll} \text{Maximize} & (\theta - \theta_k)^\top \nabla_{\theta} L_{\theta_k}(\theta)|_{\theta=\theta_k} \\ \text{subject to} & \frac{1}{2}(\theta - \theta_k)^\top H_{\theta_k}(\theta - \theta_k) \leq \delta \end{array}$$

Hessian of $\bar{D}_{KL}(\pi_{\theta_k} || \pi_{\theta})$


-
- ▶ The Hessian of $\bar{D}_{KL}(\pi_{\theta_k} || \pi_{\theta})$ is positive semi-definite
 - ▶ The constraint is therefore convex (and can be easily analyzed)
 - ▶ The solution is: usually called “natural policy gradient (NPG)”

$$\theta = \theta_k + \alpha H_{\theta_k}^{-1} \nabla_{\theta} L_{\theta_k}(\theta)|_{\theta=\theta_k} \quad (\text{HW2 problem})$$

A Primer for NPG: Scaled Gradient

$$\text{minimize}_{x \in \mathbb{R}^2} \quad f(x) := \frac{1}{2}(x - x^*)^\top Q(x - x^*)$$

Suppose $Q = [Q_{11}, 0; 0, Q_{22}]$ is a diagonal matrix with $0 < Q_{11} \ll Q_{22}$




Idea: Accelerate gradient descent by ***scaling the gradient***

$$x_{t+1} = x_t - \eta_t Q^{-1} \nabla f(x_t) = x_t - \eta_t (x_t - x^*)$$

Natural Policy Gradient (NPG)

NPG: Use **Fisher information matrix** to scale the gradient

$$\theta_{k+1} \leftarrow \theta_k + \eta \cdot \left. H_{\theta_k}^{-1} \nabla_{\theta} V^{\pi_{\theta}}(\mu) \right|_{\theta=\theta_k}$$
$$H_{\theta} := \mathbb{E}_{s \sim d_{\mu}^{\pi_{\theta}}} \left[\left(\nabla_{\theta} \log \pi_{\theta}(a|s) \right) \left(\nabla_{\theta} \log \pi_{\theta}(a|s) \right)^{\top} \right] \quad (\text{Fisher information matrix})$$



Distance metric of θ does not fully capture
“distance between probability distributions”

(For ease of presentation, assume F_{θ} has full rank; otherwise use pseudo inverse)

Quick Summary: What is TRPO?

TRPO = TR Method on RL

With 3 key steps...

1. Approximate $d_{\mu}^{\pi_{new}}(s)$ by $d_{\mu}^{\pi_{old}}(s)$
2. Trust region by KL divergence between π_{old}, π_{new}
3. Simplify the problem by linear and quadratic approximation

Assignment for this lecture:

- ▶ Spend 30 minutes going through the idea of TRPO again
- ▶ Spend 30 minutes **reading the code of TRPO**
 - ▶ <https://github.com/ikostrikov/pytorch-trpo>

- ▶ Could you explain the purpose of each line?
- ▶ Could you find any part of the code that we have not discussed in this lecture?

We will discuss this next time!

```
51 def trpo_step(model, get_loss, get_kl, max_kl, damping):
52     loss = get_loss()
53     grads = torch.autograd.grad(loss, model.parameters())
54     loss_grad = torch.cat([grad.view(-1) for grad in grads]).data
55
56     def Fvp(v):
57         kl = get_kl()
58         kl = kl.mean()
59
60         grads = torch.autograd.grad(kl, model.parameters(), create_graph=True)
61         flat_grad_kl = torch.cat([grad.view(-1) for grad in grads])
62
63         kl_v = (flat_grad_kl * Variable(v)).sum()
64         grads = torch.autograd.grad(kl_v, model.parameters())
65         flat_grad_grad_kl = torch.cat([grad.contiguous().view(-1) for grad in grads]).data
66
67         return flat_grad_grad_kl + v * damping
68
69     stepdir = conjugate_gradients(Fvp, -loss_grad, 10)
70
71     shs = 0.5 * (stepdir * Fvp(stepdir)).sum(0, keepdim=True)
72
73     lm = torch.sqrt(shs / max_kl)
74     fullstep = stepdir / lm[0]
75
76     neggdotstepdir = (-loss_grad * stepdir).sum(0, keepdim=True)
77     print(("lagrange multiplier:", lm[0], "grad_norm:", loss_grad.norm()))
78
79     prev_params = get_flat_params_from(model)
80     success, new_params = linesearch(model, get_loss, prev_params, fullstep,
81                                     neggdotstepdir / lm[0])
82     set_flat_params_to(model, new_params)
83
84     return loss
```