

535514: Reinforcement Learning

Lecture 22 — DQN, DDQN, and Distributional RL

Ping-Chun Hsieh

May 9, 2024

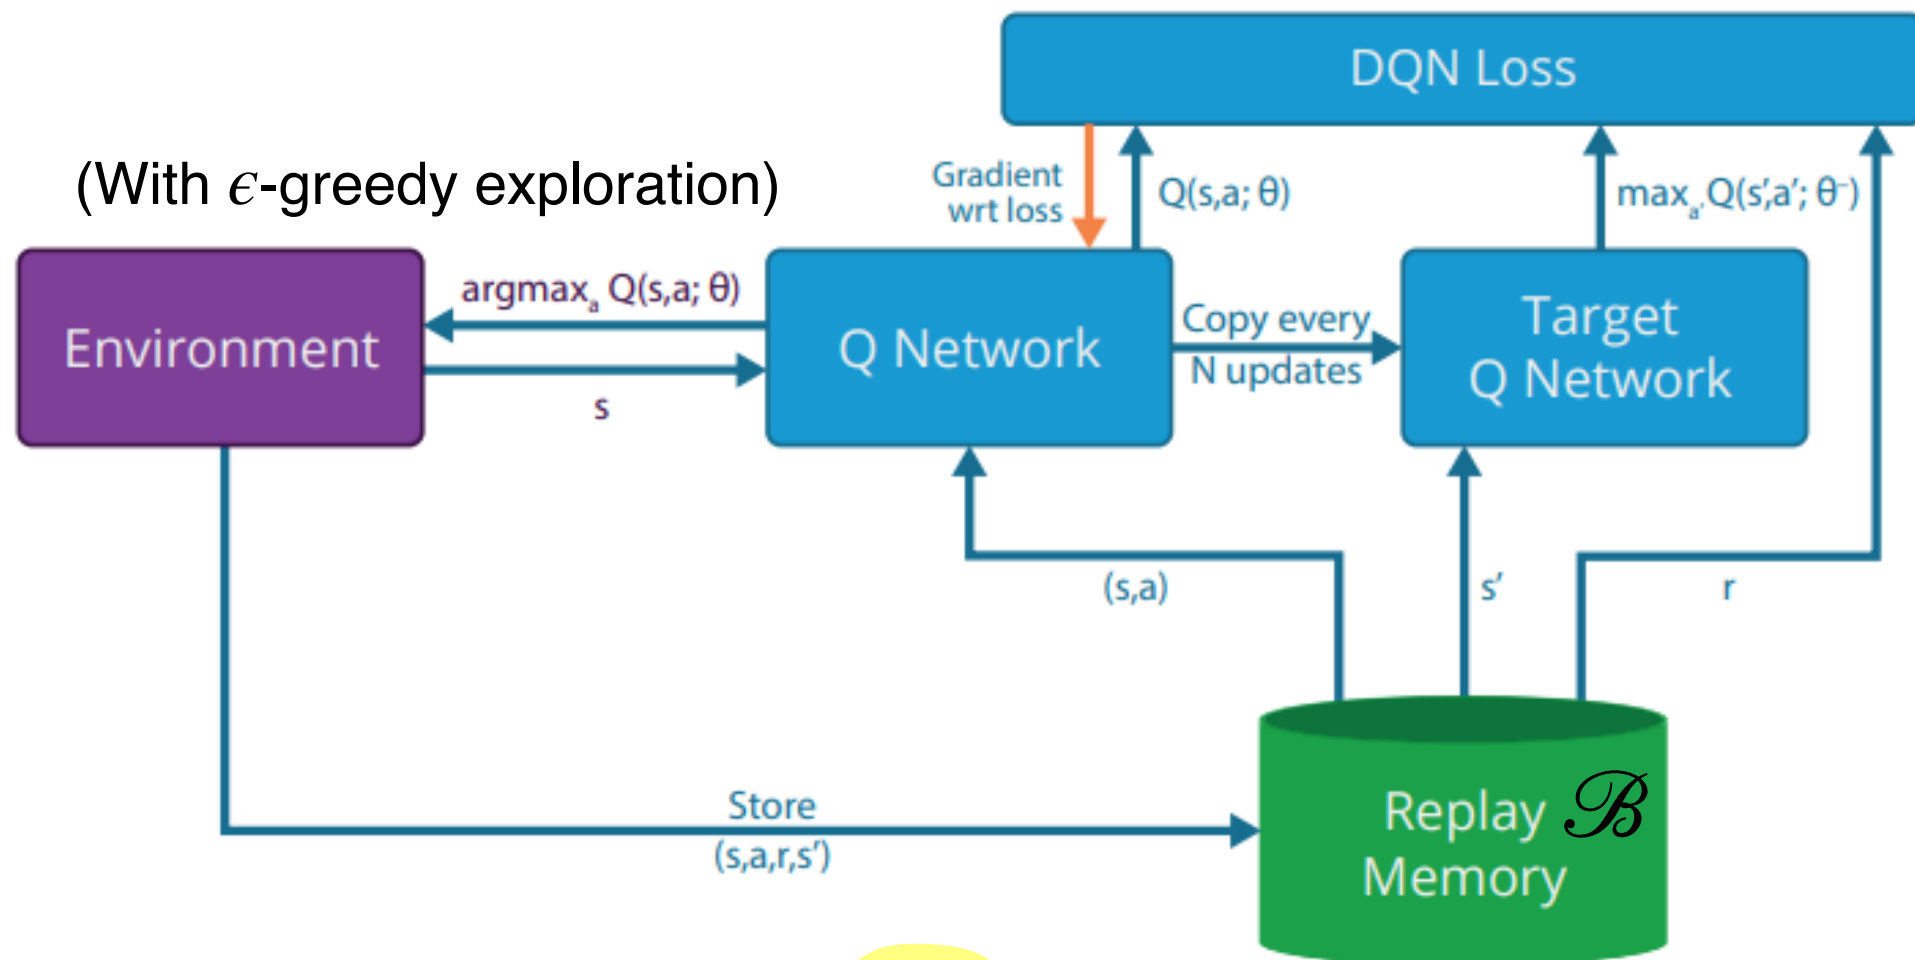
Announcements

- ▶ Team Project Milestones:
 - ▶ 1st Team-Mentor Meetup: 5/7-5/10 (Week 12)
 - ▶ 2nd Team-Mentor Meetup: 5/27-5/29 (Week 15)
 - ▶ Poster/Oral presentations: 6/11-6/13 (2.5-hour sessions, TBD)
 - ▶ Submission of technical report: by 6/17
- ▶ Theory project:
 - ▶ Submit your Hackmd note: by 5/21 (Tuesday), 9pm
 - ▶ Hackmd template: <https://hackmd.io/@pinghsieh/r1biYBHz0/edit>
 - ▶ Peer reviews: 5/22-5/28

On-Policy vs Off-Policy Methods

	Policy Optimization	Value-Based	Model-Based	Imitation-Based
On-Policy	Exact PG REINFORCE (w/i baseline) A2C On-policy DAC TRPO Natural PG (NPG) PPO-KL & PPO-Clip RLHF by PPO-KL	Epsilon-Greedy MC Sarsa Expected Sarsa	Model-Predictive Control (MPC) PETS	IRL GAIL IQ-Learn
Off-Policy	Off-policy DPG & DDPG Twin Delayed DDPG (TD3)	Q-learning Double Q-learning DQN & DDQN C51 / QR-DQN / IQN Rainbow Soft Actor-Critic (SAC)		

Review: Deep Q-Network



$$F(\mathbf{w}) := \frac{1}{2} \mathbb{E}_{(s,a,r,s') \sim \rho} \left[\left(r + \gamma \max_{a' \in A} \underbrace{Q(s', a'; \bar{\mathbf{w}})}_{\text{target network}} - Q(s, a; \mathbf{w}) \right)^2 \right]$$

$$\approx \frac{1}{2} \sum_{(s,a,r,s') \in D \sim \mathcal{B}} \left[\left(r + \gamma \max_{a' \in A} Q(s', a'; \bar{\mathbf{w}}) - Q(s, a; \mathbf{w}) \right)^2 \right]$$

(experience) ₄ relay buffer

3. Learn a generator of actions \rightarrow


Draw $a^{(1)}, \dots, a^{(N)}$ actions

Take $\arg\max_i Q(s, a^{(i)}; w)$


Can DQN be Applied Under Continuous Actions?

The difficulty lies in the “ $\arg\max_{a \in \mathcal{A}} Q(s, a; w)$ ” operation

1. Discretization




2. Gradient ascent



Existing methods that adapts DQN to continuous actions

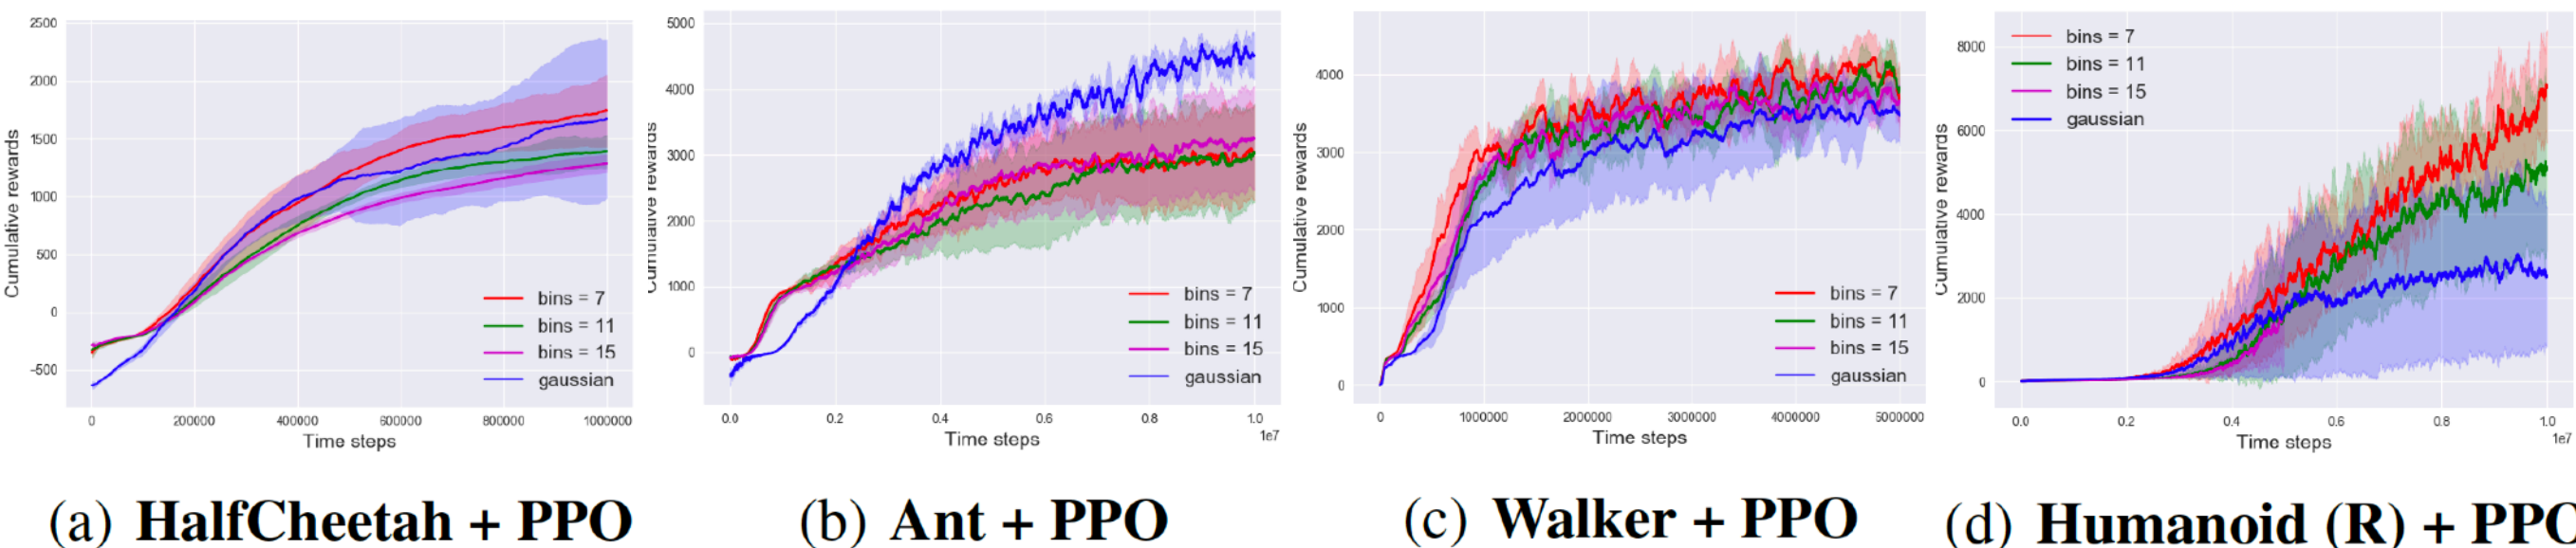
1. Naive Action Discretization [Tavakoli et al., AAAI 2018]



Issue: Naive discretization suffers from exponential growth of cardinality

2. Discretization + Factorization [Tang and Agrawal, AAAI 2020]




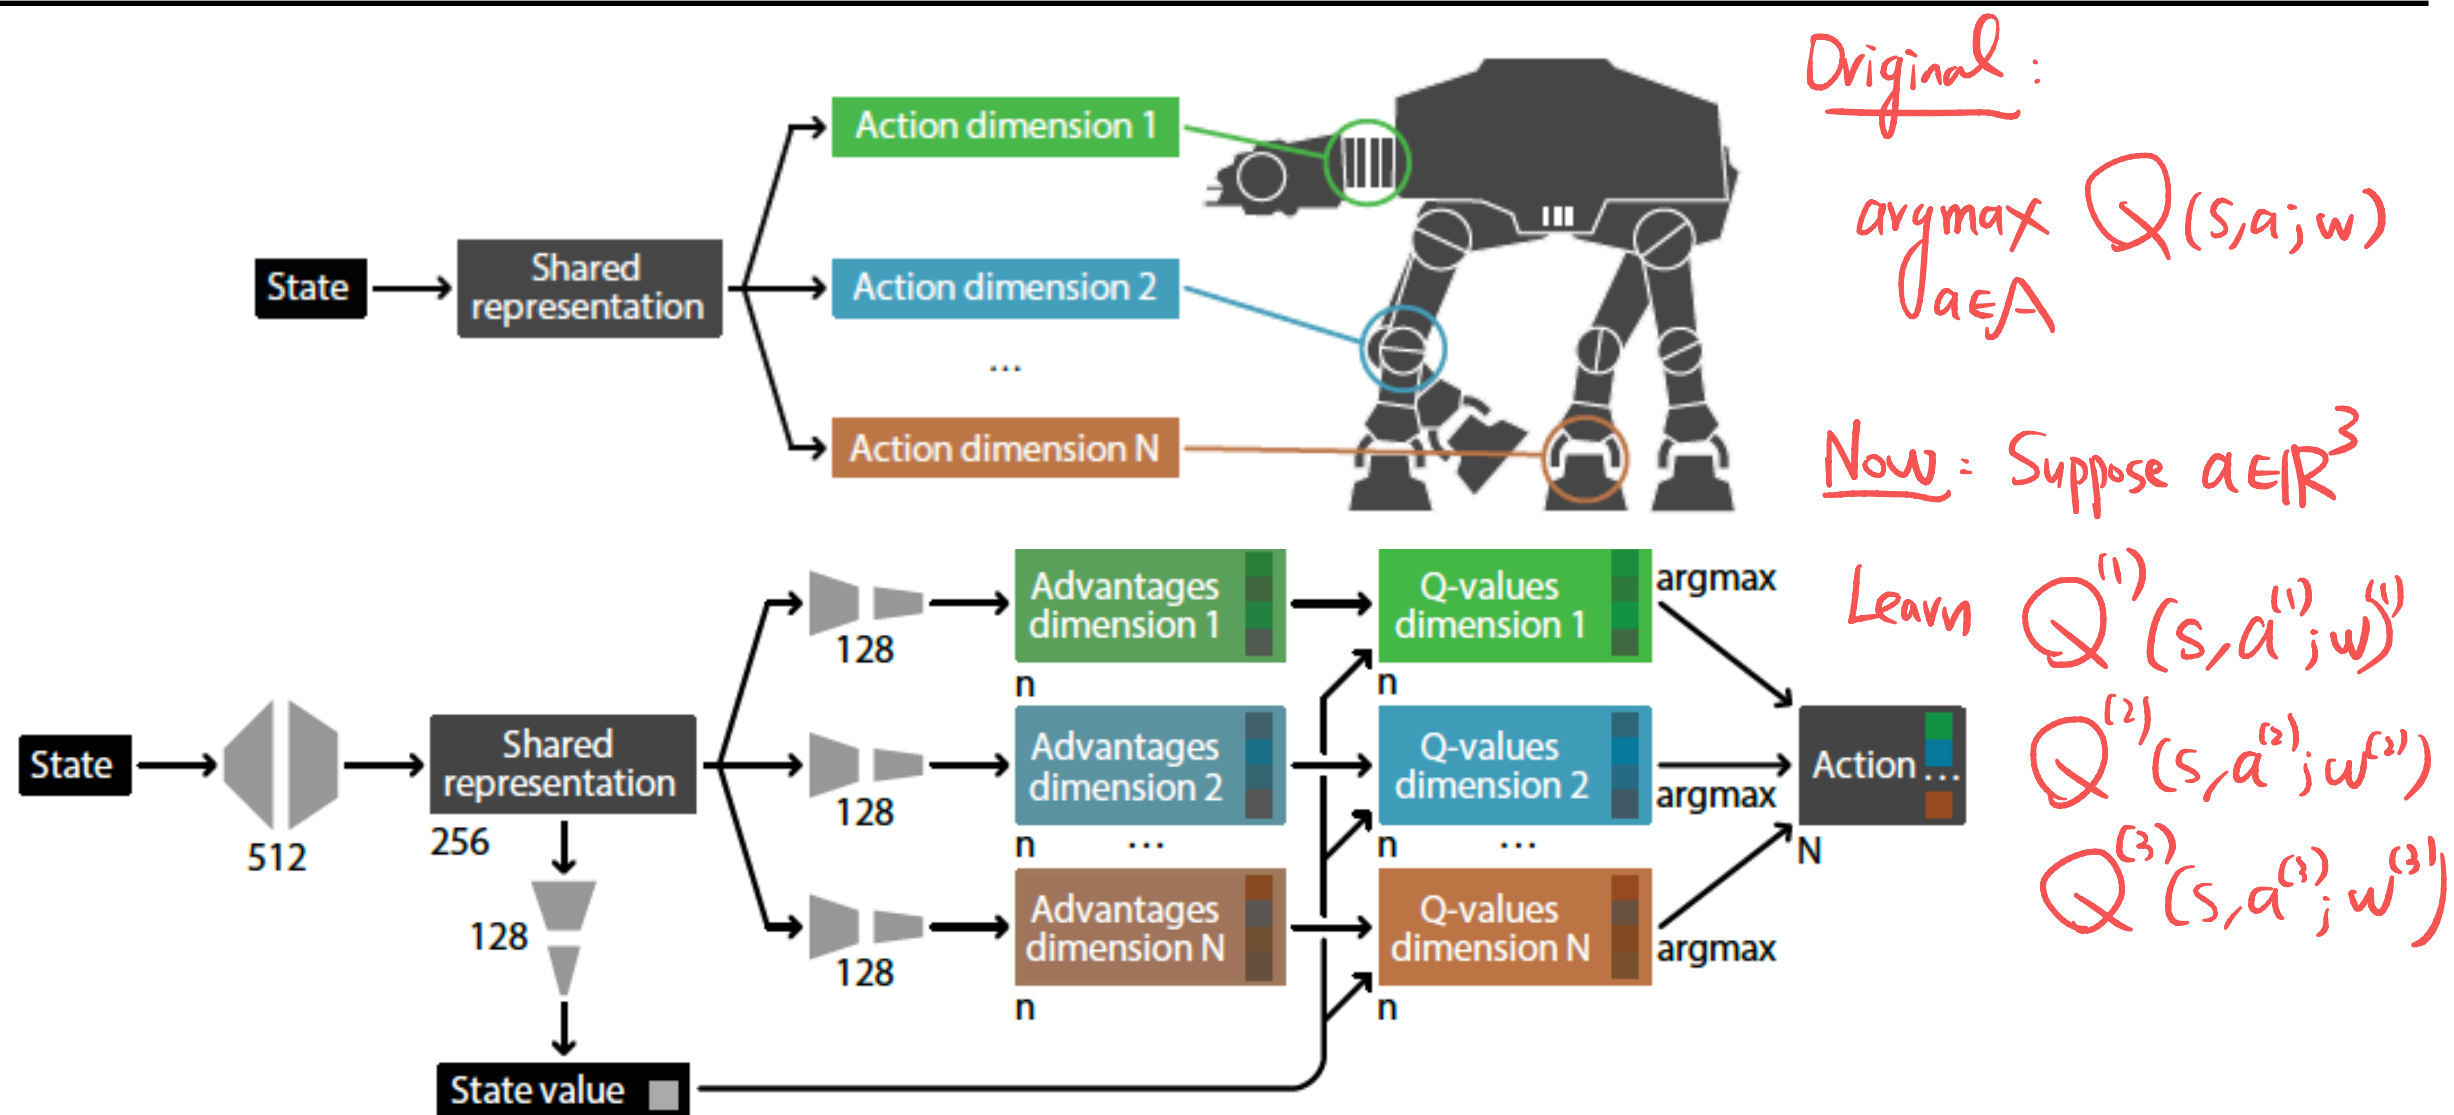
$$\pi(a|s) := \prod_{i=1}^d \pi_{\theta_i}(a_i|s) \quad \text{where } a = [a_0, a_1, \dots, a_{d-1}]^T$$



Tavakoli et al., Action Branching Architectures for Deep Reinforcement Learning, AAAI 2018

Tang and Agrawal et al., Discretizing Continuous Action Space for On-Policy Optimization, AAAI 2020

3. Discretization + Branching [Tavakoli et al., AAAI 2018]



4. Normalized Advantage Functions (NAF): Quadratic Approximation!

Continuous Deep Q-Learning with Model-based Acceleration

[Gu et al., ICML 2016]

Shixiang Gu^{1 2 3}
Timothy Lillicrap⁴
Ilya Sutskever³
Sergey Levine³

SG717@CAM.AC.UK
COUNTZERO@GOOGLE.COM
ILYASU@GOOGLE.COM
SLEVINE@GOOGLE.COM


¹University of Cambridge ²Max Planck Institute for Intelligent Systems ³Google Brain ⁴Google DeepMind

$$Q(s, a; \phi_A, \phi_V) = A(s, a; \phi_A) + V(s; \phi_V) \quad (P \text{ is state-dependent, positive definite})$$

$$A(s, a; \phi_A) := -\frac{1}{2} (a - \mu(s; \phi_\mu))^T P(s; \phi_P) (a - \mu(s; \phi_\mu)) \Rightarrow \text{The maximizer of } A(s, a; \phi_A) \text{ is simply } \mu(s; \phi_\mu)$$

$$P(s; \phi_P) := L(s; \phi_P) L(s; \phi_P)^T \quad (\text{This is known as the "Cholesky decomposition"})$$

(L is a lower-triangular matrix)



5. Amortized Q-Learning (AQL): Sampling!

Q-LEARNING IN ENORMOUS ACTION SPACES VIA AMORTIZED APPROXIMATE MAXIMIZATION

[NeurIPS 2018 Workshop]

Tom Van de Wiele*, David Warde-Farley, Andriy Mnih & Volodymyr Mnih

DeepMind


London, United Kingdom

tvdwiele@gmail.com, {dwf, amnih, vmnih}@google.com

$$\max_{a \in \mathcal{A}} Q(s, a; \mathbf{w}) \approx \max_{a' \in D \sim \mu(a)} Q(s, a'; \mathbf{w})$$

“proposal distribution”
(Learned by maximum likelihood)

Intuition: Larger $|D|$ induces a higher probability of seeing max-Q actions



$$A := \{a^{(1)}, a^{(2)}, \dots, a^{(10000)}\}$$

Suppose $a^{(i)} = \underset{a \in A}{\operatorname{argmax}} Q(s, a)$



Sampling distribution μ

Suppose $\mu(a^{(i)}) = \underline{0.01}$ and we draw K actions independently from μ .

$$P(\text{sample } a^{(i)} \text{ for at least once}) = 1 - (0.99)^{\overset{K=100}{K}}$$

6. DDPG: Reinterpret DDPG as an Adaptation of DQN for Continuous Actions!

(Quick Review)

Off-Policy Deterministic PG: $\nabla_{\theta} J_{\beta}^{\pi_{\theta}} \approx \mathbb{E}_{s \sim d_{\mu}^{\beta}} \left[\nabla_{\theta} \pi_{\theta}(s) \nabla_a Q^{\pi_{\theta}}(s, a) |_{a=\pi_{\theta}(s)} \right]$

- **Critic**: estimate $Q_w \approx Q^{\pi_{\theta}}$ by bootstrapping
- **Actor**: updates policy parameters θ by deterministic policy gradient

Step 1: Initialize θ_0 , w_0 and step sizes α_{θ} , α_w

Step 2: Sample a trajectory $\tau = (s_0, a_0, r_1, \dots) \sim P_{\mu}^{\beta}$

For each step of the current trajectory $t = 0, 1, 2, \dots$

$$\Delta w_k \leftarrow \Delta w_k + \alpha_w (r_t + \gamma Q_{w_k}(s_{t+1}, \pi_{\theta}(s_{t+1})) - Q_{w_k}(s_t, a_t)) \nabla_w Q_w(s_t, a_t) |_{w=w_k}$$

$$\Delta \theta_k \leftarrow \Delta \theta_k + \alpha_{\theta} \gamma^t \left(\nabla_{\theta} \pi_{\theta}(s_t) \nabla_a Q_{w_k}(s_t, a) |_{a=\pi_{\theta}(s_t)} \right)$$

$= \nabla_{\theta} Q_{w_k}(s_t, \pi_{\theta}(s_t)) |_{\theta=\theta_k}$

Alternative Interpretation of DDPG: An Adaptation of DQN for Continuous Actions (Cont.)

- ▶ DDPG can be reinterpreted as DQN for continuous actions

1. Deterministic policy: $\pi_{\theta}(s) \approx \arg \max_a Q_w(s, a)$

2. How to find θ : solve $\theta \leftarrow \arg \max_{\theta} Q_w(s, \pi_{\theta}(s))$ by SGD

(Actor)

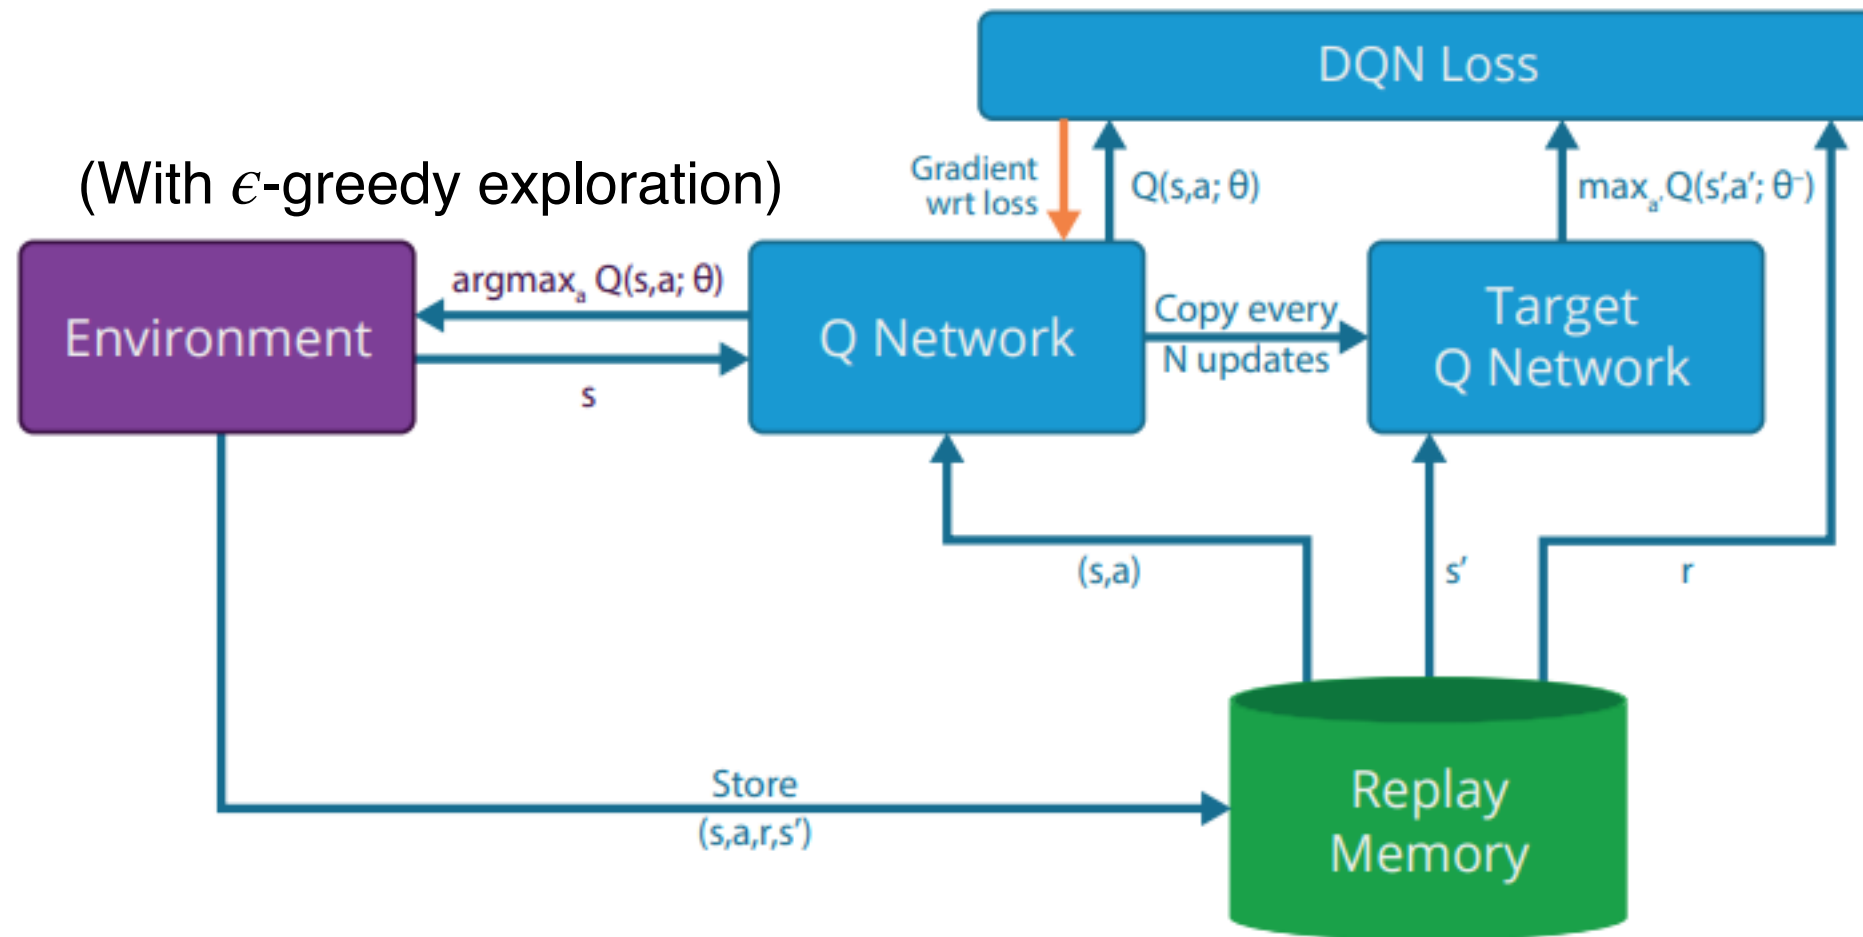
$$\Delta \theta_k \leftarrow \Delta \theta_k + \alpha_{\theta} \gamma^t \left(\nabla_{\theta} \pi_{\theta}(s_t) \nabla_a Q_{w_k}(s_t, a) \Big|_{a=\pi_{\theta}(s_t)} \right)$$

$\xrightarrow{\quad} = \nabla_{\theta} Q_{w_k}(s_t, \pi_{\theta}(s_t)) \Big|_{\theta=\theta_k}$

3. DQN and DDPG have a similar TD update scheme

$$\Delta w_k \leftarrow \Delta w_k + \alpha_w \left(r_t + \gamma Q_{w_k}(s_{t+1}, \pi_{\theta}(s_{t+1})) - Q_{w_k}(s_t, a_t) \right) \nabla_w Q_w(s_t, a_t) \Big|_{w=w_k}$$

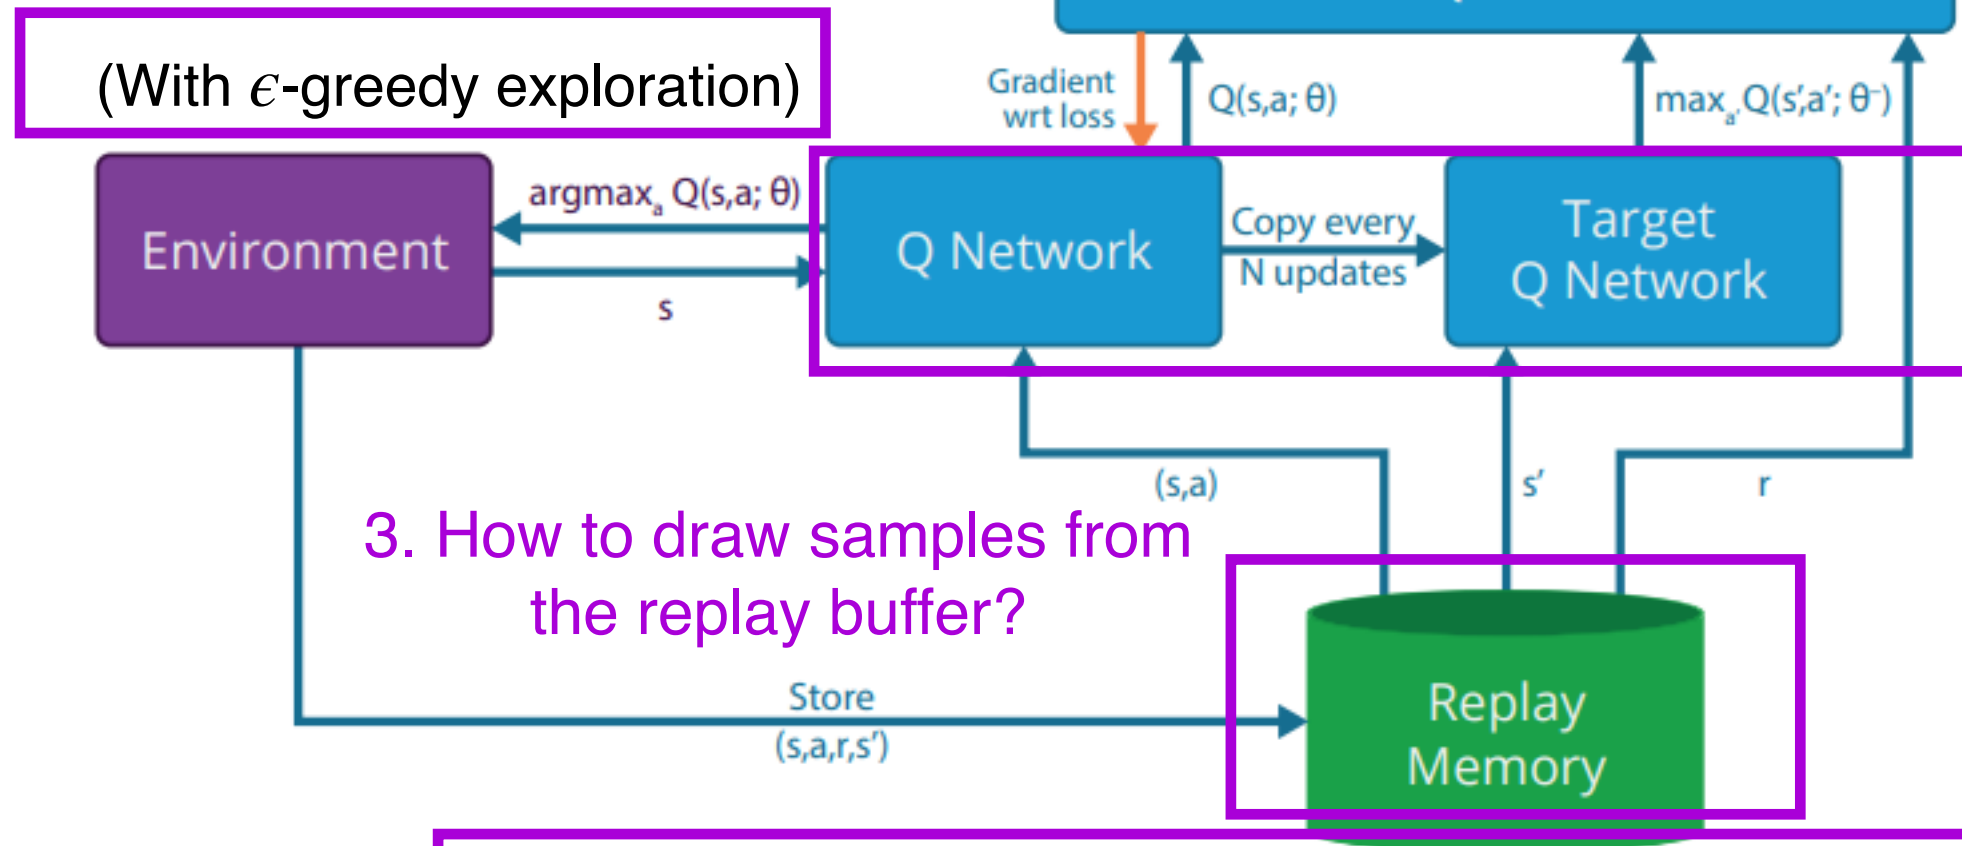
Next Topic: What to Improve in Vanilla DQN?



$$L(\mathbf{w}) := \sum_{(s,a,r,s') \in D} \frac{1}{2} \left[\left(r + \gamma \max_{a'} Q(s', a'; \bar{\mathbf{w}}) - Q(s, a; \mathbf{w}) \right)^2 \right]$$

Next Topic: What to Improve in Vanilla DQN?

4. A better exploration method?



3. How to draw samples from the replay buffer?

$$L(\mathbf{w}) := \sum_{(s, a, r, s') \in D} \frac{1}{2} \left[\left(r + \gamma \max_{a'} Q(s', a'; \bar{\mathbf{w}}) - Q(s, a; \mathbf{w}) \right)^2 \right]$$


1. Overestimation Bias?

2. A better loss function?

5. A better way to represent Q function?

Next Topic: Rainbow (= DQN with 6 Useful Tricks)

1. Double DQN (DDQN)
2. Distributional Q-learning
3. Prioritized experience replay (PER)
4. Dueling networks
5. Multi-step return in TD target
6. Noisy networks for exploration



Double DQN (DDQN)

Hado van Hasselt, Arthur Guez, and David Silver,
“Deep Reinforcement Learning with Double Q-learning,” AAAI 2016

Recall: Double Q-Learning

Step 1: Initialize $Q^A(s, a)$, $Q^B(s, a)$ for all (s, a) , and initial state s_0

Step 2: For each step $t = 0, 1, 2, \dots$

Select a_t using ε -greedy w.r.t $Q^A(s_t, a) + Q^B(s_t, a)$

Observe (r_{t+1}, s_{t+1})

Choose one of the following updates uniformly at random

$$Q^A(s_t, a_t) \leftarrow Q^A(s_t, a_t) + \alpha(r_{t+1} + \gamma Q^B(s_{t+1}, \arg \max_a Q^A(s_{t+1}, a)) - Q^A(s_t, a_t))$$

$$Q^B(s_t, a_t) \leftarrow Q^B(s_t, a_t) + \alpha(r_{t+1} + \gamma Q^A(s_{t+1}, \arg \max_a Q^B(s_{t+1}, a)) - Q^B(s_t, a_t))$$

- ▶ **Key Idea:** Decouple “Q value” and “greedy action selection”
- ▶ **Question:** How to apply this to DQN?

Double DQN

- Loss function of DQN:

$$F(\mathbf{w}) := \frac{1}{2} \mathbb{E}_{(s,a,r,s') \sim \rho} \left[\left(r + \gamma \max_{a' \in A} Q(s', a'; \bar{\mathbf{w}}) - Q(s, a; \mathbf{w}) \right)^2 \right]$$
$$\approx \frac{1}{2} \sum_{(s,a,r,s') \in D} \left[\left(r + \gamma \max_{a' \in A} Q(s', a'; \mathbf{w}) - Q(s, a; \mathbf{w}) \right)^2 \right]$$

- Loss function of Double DQN:

$$F(\mathbf{w}) := \frac{1}{2} \mathbb{E}_{(s,a,r,s') \sim \rho} \left[\left(r + \gamma Q(s', \arg \max_{a' \in A} Q(s, a; \mathbf{w}); \bar{\mathbf{w}}) - Q(s, a; \mathbf{w}) \right)^2 \right]$$
$$\approx \frac{1}{2} \sum_{(s,a,r,s') \sim D} \left[\left(r + \gamma Q(s', \arg \max_{a' \in A} Q(s, a; \mathbf{w}); \bar{\mathbf{w}}) - Q(s, a; \mathbf{w}) \right)^2 \right]$$

“We therefore propose to *evaluate the greedy policy according to the online network*, but using the target network to estimate its value.” —
[van Hasselt et al., AAAI 2016]

Distributional Q-Learning

(Learn value distribution $Z(s, a)$ & use $E[Z(s, a)]$ as $Q(s, a)$ in Q-Learning)

Why Shall We Consider “Value Distributions”?

- ▶ **Risky vs safe choices**
 - ▶ E.g., Same expected return but different variance
- ▶ **Good empirical performance** (despite that the underlying root cause is not fully known)
 - ▶ C51 [Belleware et al., ICML 2017]
 - ▶ QR-DQN [Dabney et al., AAAI 2018]
 - ▶ IQN [Dabney et al., ICML 2018]
- ▶ **New approaches for exploration**
 - ▶ Information-directed exploration [Nikolov et al., ICLR 2019]
 - ▶ Distributional RL for efficient exploration [Mavrin et al., ICML 2019]
- ▶ **Learn better critics**
 - ▶ Truncated Quantile Critics (TQC) [Kuznetsov et al., ICML 2020]


Question: How to learn the complete value distribution (instead of merely the expectation)?

Sample Action-Value $Z^\pi(s, a)$

- ▶ Sample action-value $Z^\pi(s, a)$: sample return if we start from state s and take action a , and then follow policy π

$$Q^\pi(s, a) = \mathbb{E}[Z^\pi(s, a)] = \mathbb{E}\left[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t)\right]$$

- ▶ $Z^\pi(s, a)$ is essentially a random variable (and hence follows some distribution)
- ▶ **Example**: 1-state MDP with 2 actions and $\pi(s) = a_1$
 - ▶ $Q^\pi(s, a_1) = ?$ $Z^\pi(s, a_1) = ?$



- ▶ $Q^\pi(s, a_2) = ?$ $Z^\pi(s, a_2) = ?$

Finding Z^π via Distributional Bellman Equation

- ▶ **Mild assumption:** $Z^\pi(s, a)$ has bounded moments
- ▶ Distributional Bellman equation for $Z^\pi(s, a)$: Given s, a , we have

$$Z^\pi(s, a) \stackrel{D}{=} r(s, a) + \gamma Z^\pi(s', a')$$

($\stackrel{D}{=}$: equal in distribution)

- ▶ **Question:** How to interpret this equation?
- ▶ **Question:** Are $r(s, a)$ and $Z^\pi(s', a')$ independent?
- ▶ **Question:** Is this consistent with Bellman expectation equation?

Distributional Bellman Operator B^π

- ▶ \mathcal{Z} : the space of all value distributions with bounded moments
- ▶ Transition operator $P^\pi : \mathcal{Z} \rightarrow \mathcal{Z}$

$$P^\pi Z(s, a) \stackrel{D}{=} Z(s', a')$$


$$s' \sim P(\cdot | s, a), \quad a' \sim \pi(\cdot | s')$$

- ▶ Distributional Bellman operator $B^\pi : \mathcal{Z} \rightarrow \mathcal{Z}$

$$B^\pi Z(s, a) \stackrel{D}{=} r(s, a) + \gamma P^\pi Z(s, a)$$

An Example of Applying B^π

- ▶ **Example:** 2 states x_1, x_2 and 2 actions a_1, a_2
- ▶ $\pi(a_1 | x_2) = 0.3$, $\pi(a_2 | x_2) = 0.7$, and $\gamma = 0.9$






$$B^\pi Z(s, a) \stackrel{D}{=} r(s, a) + \gamma P^\pi Z(s, a)$$

- ▶ Suppose $Z(x_1, a_1) = 0$, $Z(x_2, a_1) = 0$ with probability 1 and $Z(x_2, a_2) \sim \mathcal{N}(0, 1)$
- ▶ **Question:** $B^\pi Z(x_2, a_2) = ?$ $B^\pi Z(x_1, a_1) = ?$

Visualization of Distributional Bellman Operator

$$P^\pi Z$$




Distributional “Optimality” Operator

Recall— Distributional Bellman operator $B^\pi : \mathcal{Z} \rightarrow \mathcal{Z}$

$$B^\pi Z(s, a) \stackrel{D}{=} r(s, a) + \gamma P^\pi Z(s, a)$$

- Distributional optimality operator B^* : The B^π resulting from a greedy policy π (what does “greedy” mean here?)

An Example of B^*



Suppose we have the following:

- ▶ $\pi(a_1 | x_2) = 0.3$, $\pi(a_2 | x_2) = 0.7$, and $\gamma = 1$
- ▶ $Z(x_1, a_1) = 0$, $Z(x_2, a_1) = 0$ with probability 1
- ▶ $Z(x_2, a_2) \sim \mathcal{N}(0, 1)$

Question: What's the PDF of $B^*Z(x_1, a_1) = ?$

$$B^*Z(s, a) \stackrel{D}{=} r(s, a) + \gamma P^{\pi_{\text{greedy}}} Z(s, a)$$