

535514: Reinforcement Learning

Lecture 24 — QR-DQN and IQN

Ping-Chun Hsieh

May 16, 2024

Announcement

- No class next Monday (5/20) and next Thursday (5/23)

5月

日 一 二 三 四 五 六

1 2 3 4
5 6 7 8 9 10 11

12 13 14 15 16 17 18
No class No class

19 20 21 22 23 24 25

26 27 28 29 30 31

Lec 25 (SAC)

Lec 26
(Inverse RL)

6月

日 一 二 三 四 五 六

Lec 27 (Model-based RL) Lec 28 (Offline MBRL)

2 3 4 5 6 7 8
Final presentation

9 10 11 12 13 14 15

16 17 18 19 20 21 22

23 24 25 26 27 28 29

30

On-Policy vs Off-Policy Methods

	Policy Optimization	Value-Based	Model-Based	Imitation-Based
On-Policy	Exact PG REINFORCE (w/i baseline) A2C On-policy DAC TRPO Natural PG (NPG) PPO-KL & PPO-Clip RLHF by PPO-KL	Epsilon-Greedy MC Sarsa Expected Sarsa	Model-Predictive Control (MPC) PETS	IRL GAIL IQ-Learn
Off-Policy	Off-policy DPG & DDPG Twin Delayed DDPG (TD3)	Q-learning Double Q-learning DQN & DDQN Rainbow C51 / QR-DQN / IQN Soft Actor-Critic (SAC)		

Quick Review: Distributional Bellman

- ▶ Sample action-value $Z^\pi(s, a)$: sample return if we start from state s and take action a , and then follow policy π

$$\underline{Q}^\pi(s, a) = \mathbb{E}[\underline{Z}^\pi(s, a)] = \mathbb{E}\left[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t)\right]$$

- ▶ Distributional Bellman operator $B^\pi : \mathcal{Z} \rightarrow \mathcal{Z}$

$$B^\pi Z(s, a) \stackrel{D}{=} \underline{r(s, a)} + \gamma \underline{P^\pi Z(s, a)}$$

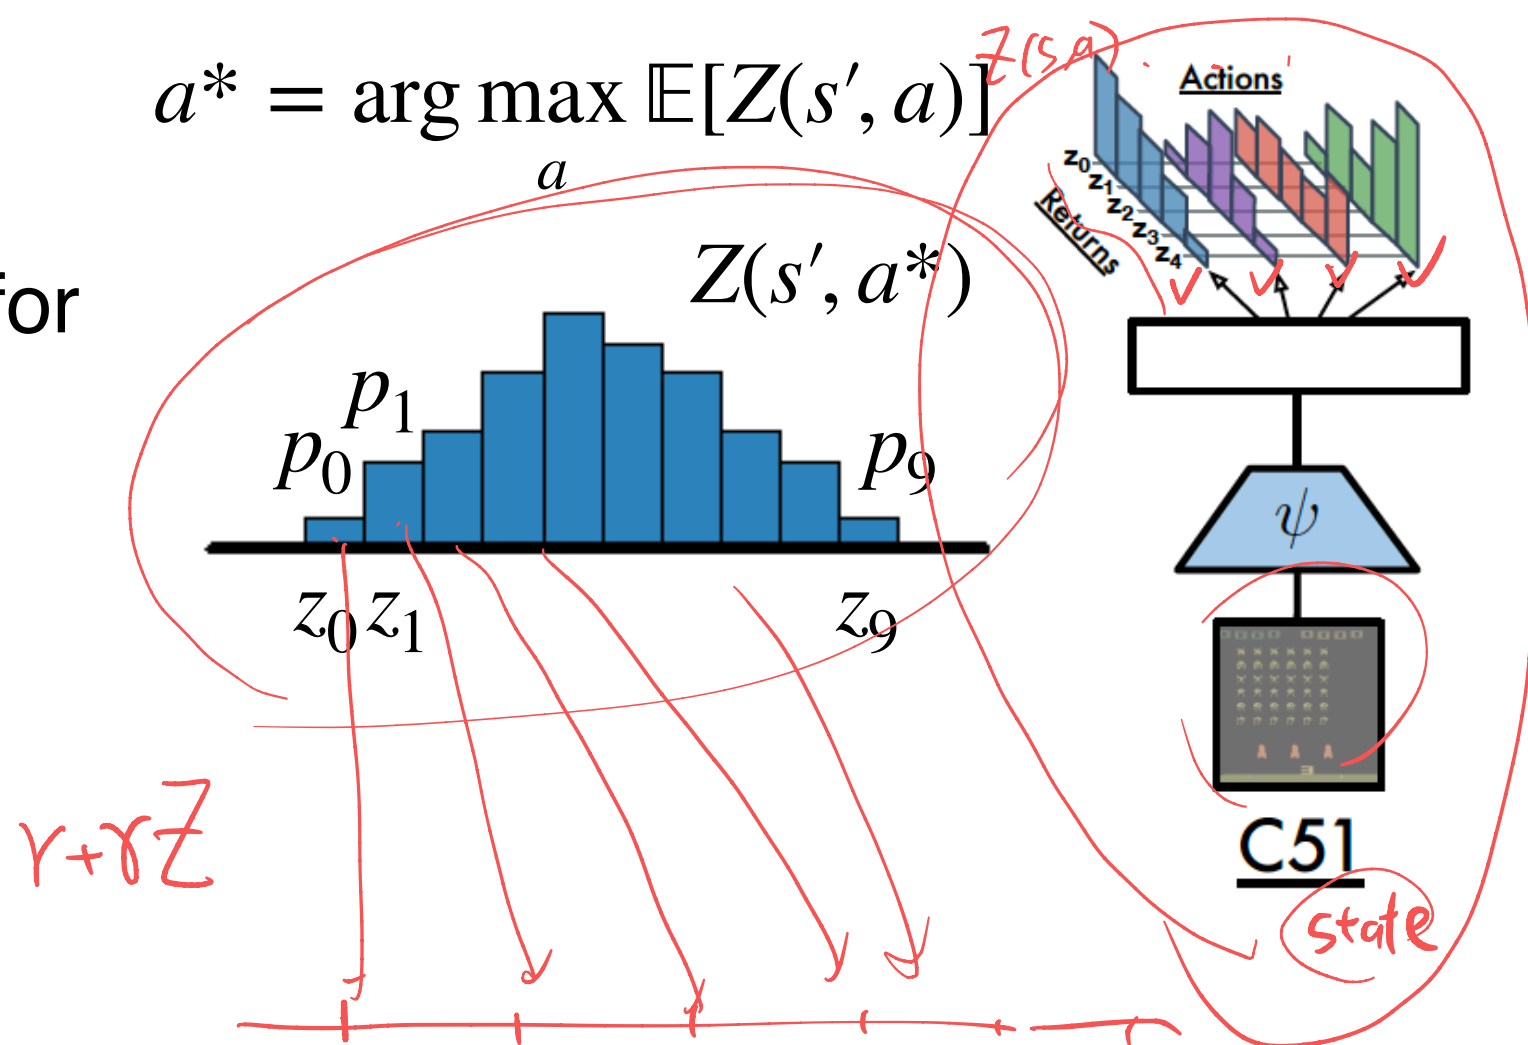
where $P^\pi Z(s, a) \stackrel{D}{=} Z(s', a')$
 $s' \sim P(\cdot | s, a), a' \sim \pi(\cdot | s')$

- ▶ Distributional optimality operator B^* : The B^π resulting from a greedy policy π (what does “greedy” mean here?)

Quick Review: C51

(C1) Categorical distributions for parametrizing $Z_\theta(s, a)$

V_{max} , V_{min}



(C2) Mimicking B^* for learning with sample transitions (s, a, r, s')


(C3) Cramer Projection Φ for support mismatch caused by $B^*Z_\theta(s, a)$

(C4) Minimize $L_{C51}(s, a, r, s'; \theta) := D_{KL}(\Phi B^*Z_{\bar{\theta}}(s, a) || Z_\theta(s, a))$

QR-DQN

Quantile-Based Parametrization of $Z(s, a)$

- **Idea:** Express $Z(s, a)$ using CDF (instead PDF)




(inverse CDF)


Quantile function: $F_Z^{-1}(\tau) := \inf\{z : P(Z \leq z) \geq \tau\}$

$$P(Z(s, a) \leq v_1) = \tau_1$$

PMF ($P(Z(s, a) = v_i)$)



A Comparison of NN Architecture



QR-DQN: Another Popular Distributional DQN

✓ Q1: How to express $Z(s, a)$?


(D1) **Quantile** distributions for $Z_\theta(s, a)$

✓ Q2: How to update $Z(s, a)$ during training?

(D2) Mimicking B^* for learning with sample transitions (s, a, r, s')

(D3) Minimize $L_{QR}(s, a, r, s'; \theta) := D(\underbrace{B^* Z_{\bar{\theta}}(s, a)}_{\text{No Cramer projection required!}} \| Z_\theta(s, a))$


No Cramer projection required!



Quantile Regression DQN (Formally)

Step 1: Initialize $Z_\theta(s, a)$ and initial state s_0

quantile distribution



Step 2: For each step $t = 0, 1, 2, \dots$

Select a_t using ε -greedy w.r.t $Q(s_t, a) \equiv \mathbb{E}[Z_\theta(s_t, a)]$

Observe (r_{t+1}, s_{t+1}) and store $(s_t, a_t, r_{t+1}, s_{t+1})$ in the buffer

Draw a mini-batch of samples B from the replay buffer

Update θ by minimizing QR loss as follows:

$$\theta \leftarrow \theta - \alpha \nabla_\theta \sum_{(s, a, r, s') \in B} L_{QRDQN}(s, a, r, s'; \theta)$$

Under quantile distributions, $\mathbb{E}[Z_\theta(s, a)] = \sum_{i=1}^N \frac{1}{N} Z_\theta^{(i)}(s, a)$


(D2) Mimicking B^* for Learning With Sample Transitions

$$B^* Z(s, a) = Y(s, a) + \gamma \underbrace{P^{\pi_{\text{greedy}}}}_{\text{greedy policy}} Z$$

- Here we presume a greedy policy w.r.t Q function for B^* $Z(s', a^*)$
- Question:** Given only transitions (s, a, r, s') , how to enforce B^* to update $Z(s, a)$ on *quantile* distributions?

$$a^* = \arg \max_a Q(s', a) \equiv \arg \max_a \mathbb{E}[Z_{\bar{\theta}}(s', a)]$$


$Z_{\bar{\theta}}(s', a^*)$



$$B^* Z_{\bar{\theta}}(s, a) = r + \gamma Z_{\bar{\theta}}(s', a^*)$$

Suppose $\gamma = 0.9$

$r = 1.5$



(D3) Loss Function

- ▶ We still need to choose a “**dissimilarity**” function $D(\cdot \| \cdot)$ in $L_{QRDQN}(s, a, r, s'; \theta) := D(B^*Z_{\bar{\theta}}(s, a) \| Z_{\theta}(s, a))$
- ▶ There are many possibilities, e.g., total variation or KL divergence
- ▶ QR-DQN uses the **quantile regression loss**
 - ▶ Motivation: Both $B^*Z_{\bar{\theta}}(s, a)$ and $Z_{\theta}(s, a)$ are quantile distributions


Quantile Regression Loss

- Given a random variable Z
- Goal: Find a quantile $F_Z^{-1}(\tau)$

- Idea: Finding a quantile $F_Z^{-1}(\tau)$ by minimizing loss $L(x; Z, \tau)$

$$F_Z^{-1}(\tau) = \arg \min_{x \in \mathbb{R}} L(x; Z, \tau)$$

- $L(x; Z, \tau)$ is *easy-to-optimize* when it is **strictly convex**



The Quantile Regression Loss

- ▶ Given that the derivative of $L(x; Z, \tau)$ is $F_Z(x) - \tau$, we can recover the QR loss by integration

Quantile regression (QR) loss:

$$L_{QR}(x; Z, \tau) = (\tau - 1) \int_{-\infty}^x (z - x) dF_Z(z) + \tau \int_x^{\infty} (z - x) dF_Z(z)$$

(It is easy to verify that $\frac{d}{dx} L_{QR}(x; Z, \tau) = F_Z(x) - \tau$ by the Leibniz integral rule)

Alternative expression of QR loss:

$$\begin{aligned} \frac{d}{dx} \left(\int_{a(x)}^{b(x)} f(x, t) dt \right) \\ = f(x, b(x)) \cdot \frac{d}{dx} b(x) - f(x, a(x)) \cdot \frac{d}{dx} a(x) + \int_{a(x)}^{b(x)} \frac{\partial}{\partial x} f(x, t) dt \end{aligned}$$

$$\rho_{\tau}(y) := y(\tau - \mathbb{I}\{y < 0\})$$

$$L_{QR}(x; Z, \tau) = E_Z[\rho_{\tau}(Z - x)]$$

Summary: Loss Function of QR-DQN

$$L_{QRDQN}(s, a, r, s'; \theta) := \sum_{i=1}^N L_{QR}(B^*Z_{\bar{\theta}}(s, a); Z_{\theta}(s, a), \tau_i)$$

$$= \sum_{i=1}^N \mathbb{E}_{z \sim B^*Z_{\bar{\theta}}(s, a)} [\rho_{\tau_i}(z - Z_{\theta}(s, a))]$$

N different quantiles


- **Question:** Is $L_{QRDQN}(s, a, r, s'; \theta)$ easy to compute during training?

Implicit Quantile Networks (IQN)

Dabney et al., Implicit Quantile Networks for Distributional Reinforcement Learning, ICML 2018

IQN: A Generative Approach to Distributional RL

- An illustrative comparison of **distributional** Q-learning methods



Distributional RL via explicitly expressing the distribution $Z(s, a)$

Distributional RL via a **generative model** for distribution $Z(s, a)$

➡ Need sufficiently **many atoms or quantiles** for an accurate representation of $Z(s, a)$

Calculate QR Loss by *Sampling*

QR loss:


$$\rho_{\tau}(y) := y(\tau - \mathbb{I}\{y < 0\})$$

$$L_{QR}(x; Z, \tau) = E_{z \sim Z}[\rho_{\tau}(z - x)]$$

- ▶ Recall QR-DQN:
 - ▶ The QR loss is calculated **explicitly**
 - ▶ $Z \Rightarrow$ target distribution induced by $\{\bar{\theta}_1, \dots, \bar{\theta}_N\}$
- ▶ **Idea**: A practical way to calculate the QR loss is **sampling**!

$$L_{QR}(x; Z, \tau) \approx$$

- ▶ IQN **implicitly** parameterizes Z by constructing a **generator** for Z



Suppose we are given a distribution of Z , denoted by F_Z (CDF).

Q: How to generate "random variables" of CDF F_Z ?

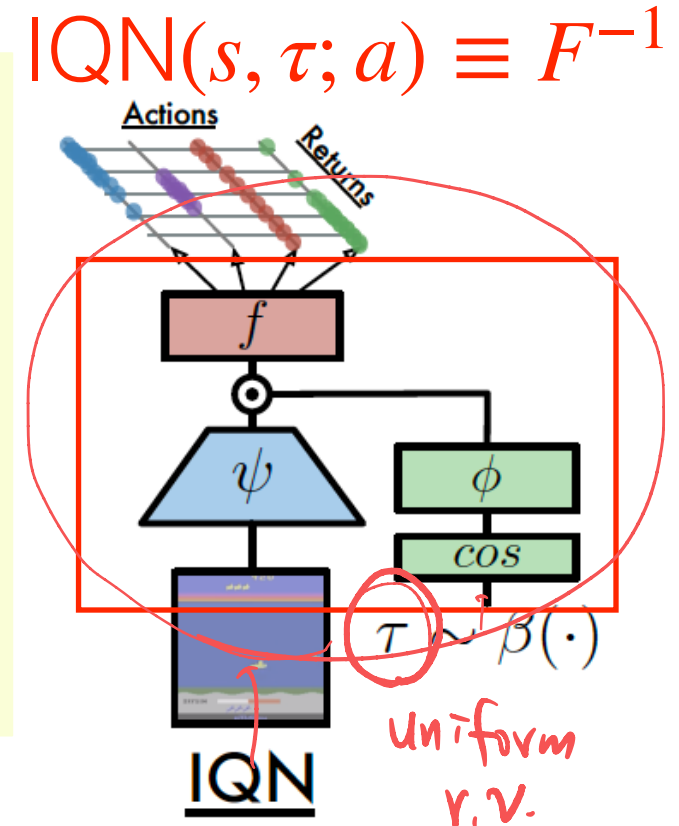
I T S
↓ ↓ ↘ sampling
inverse transform

QR Loss and Inverse Transform Sampling

QR loss:

$$\rho_\tau(y) := y(\tau - \mathbb{I}\{y < 0\})$$
$$\underline{L(x; Z, \tau)} = \underline{E_{z \sim Z}[\rho_\tau(z - x)]} \approx \frac{1}{K} \sum_{k=1}^K \rho_\tau(z_k - x)$$

$(z_1, \dots, z_K \sim Z)$



Inverse Transform Sampling (ITS): Generate any random variable with CDF F from a uniform random variable

1. Generate a random variable $U \sim \text{Unif}(0,1)$
2. Let $X = F^{-1}(U)$, where $F^{-1}(u) := \inf\{z : F(z) \geq u\}$


► ITS is essentially a **generative** approach!

Calculating QR Loss in IQN

QR loss:

$$\rho_{\tau}(y) := y(\tau - \mathbb{I}\{y < 0\})$$
$$L(x; Z, \tau) = E_{z \sim Z}[\rho_{\tau}(z - x)] \approx \frac{1}{K} \sum_{k=1}^K \rho_{\tau}(z_k - x)$$

$(z_1, \dots, z_K \sim Z)$



(Recall that Z corresponds to the target distribution in QR-DQN)

At each update, given (s, a, r, s') , for a given $\tau \in [0,1]$:

1. Draw $\tau'_1, \dots, \tau'_K \sim \text{Unif}(0,1)$ ← a generative step!

2. Get z_1, \dots, z_K by $z_i = r + \gamma \cdot \overline{\text{IQN}}(s', a'; \tau'_i)$

3. QR loss in IQN = $\frac{1}{K} \sum_{i=1}^K \rho_{\tau}(z_i - \text{IQN}(s, a; \tau))$

20 (can be readily extended to multiple τ)

IQN is closely related to the [reparameterization trick](#)

- ▶ Suppose we want to compute a loss $L(\theta) = E_{X \sim p_\theta}[f(X)]$
 - ▶ X is a random variable, and p_θ is the underlying distribution of X

- ▶ **Question:** $\nabla_\theta L(\theta) = ?$

$$\begin{aligned}\nabla_\theta L(\theta) &= \nabla_\theta E_{X \sim p_\theta}[f(X)] = \nabla_\theta \left(\int f(x) p_\theta(x) dx \right) \\ &= \int \left(f(x) \frac{1}{p_\theta(x)} \nabla_\theta p_\theta(x) \right) p_\theta(x) dx \\ &= \int \left(f(x) \underbrace{\nabla_\theta \log p_\theta(x)}_{\text{Easy to evaluate?}} \right) p_\theta(x) dx\end{aligned}$$

- ▶ **Reparameterization trick:** $\varepsilon \sim p(\varepsilon), L(\theta) = E_{\varepsilon \sim p}[g_\theta(\varepsilon)]$

$$\nabla_\theta L(\theta) = \nabla_\theta E_{\varepsilon \sim p}[g_\theta(\varepsilon)] = E_{\varepsilon \sim p}[\nabla_\theta g_\theta(\varepsilon)] \approx \frac{1}{K} \sum_{i=1}^K \nabla_\theta g_\theta(\varepsilon_i)$$

$(\varepsilon_1, \dots, \varepsilon_K \sim p)$